

# Importing Libraries

```
In [1]: import pandas as pd
import numpy as np
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: import re
import nltk
nltk.download('stopwords')
from nltk.util import pr
stemmer = nltk.SnowballStemmer("english")
from nltk.corpus import stopwords
import string
stopword = set(stopwords.words("english"))
```

[nltk\_data] Downloading package stopwords to  
[nltk\_data] C:\Users\mshiv\AppData\Roaming\nltk\_data...  
[nltk\_data] Package stopwords is already up-to-date!

# Loading Dataset

```
In [3]: df = pd.read_csv(r"E:\Data Science Projects\Project Files\Capstone 1 - Hate Speech\hate_speech.csv")
```

```
In [4]: df.head()
```

Out[4]:

	Unnamed: 0	count	hate_speech	offensive_language	neither	class	tweet
0	0	3	0		0	3	2 !!! RT @mayasolovely: As a woman you shouldn't...
1	1	3	0		3	0	!!!! RT @mleew17: boy dats cold...tyga dwn ba...
2	2	3	0		3	0	!!!!!! RT @UrKindOfBrand Dawg!!!! RT @80sbaby...
3	3	3	0		2	1	!!!!!!! RT @C_G_Anderson: @viva_based she lo...
4	4	6	0		6	0	!!!!!!!!!!!! RT @ShenikaRoberts: The shit you...

In [5]: `df.tail()`

Out[5]:

		Unnamed: 0	count	hate_speech	offensive_language	neither	class	tweet
<b>24778</b>	25291	3	0			2	1	1 you's a muthaf***in lie @#8220;@LifeAsKing: @2...
<b>24779</b>	25292	3	0			1	2	2 you've gone and broke the wrong heart baby, an...
<b>24780</b>	25294	3	0			3	0	1 young buck wanna eat!... dat nigguh like I ain...
<b>24781</b>	25295	6	0			6	0	1 youu got wild bitches tellin you lies
<b>24782</b>	25296	3	0			0	3	2 ~~Ruffled   Ntac Eileen Dahlia - Beautiful col...

## Data Preparation

In [6]: `df['labels'] = df['class'].map({0:"Hate Speech Detected",1:"Offensive language de...")`

In [7]: `df.head()`

Out[7]:

		Unnamed: 0	count	hate_speech	offensive_language	neither	class	tweet	labels
<b>0</b>	0	3	0			0	3	2 !!! RT @mayasolovely: As a woman you shouldn't...	NO hate and offensive speech
<b>1</b>	1	3	0			3	0	1 !!!! RT @mleew17: boy dats cold...tyga dwn ba...	Offensive language detected
<b>2</b>	2	3	0			3	0	1 !!!!!! RT @UrKindOfBrand Dawg!!!! RT @80sbaby...	Offensive language detected
<b>3</b>	3	3	0			2	1	1 !!!!!!! RT @C_G_Anderson: @viva_based she lo...	Offensive language detected
<b>4</b>	4	6	0			6	0	1 !!!!!!!! RT @ShenikaRoberts: The shit you...	Offensive language detected

In [8]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 24783 entries, 0 to 24782
Data columns (total 8 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Unnamed: 0        24783 non-null   int64  
 1   count            24783 non-null   int64  
 2   hate_speech      24783 non-null   int64  
 3   offensive_language 24783 non-null   int64  
 4   neither          24783 non-null   int64  
 5   class             24783 non-null   int64  
 6   tweet             24783 non-null   object  
 7   labels            24783 non-null   object  
dtypes: int64(6), object(2)
memory usage: 1.5+ MB
```

In [9]: `num_data=[x for x in df.columns if df[x].dtypes!='object']`

In [10]: `cat_data=[x for x in df.columns if df[x].dtypes=='object']`

In [11]: `num_data`

Out[11]: `['Unnamed: 0',  
 'count',  
 'hate_speech',  
 'offensive_language',  
 'neither',  
 'class']`

In [12]: `num_data.remove('Unnamed: 0')`

In [13]: `num_data`

Out[13]: `['count', 'hate_speech', 'offensive_language', 'neither', 'class']`

In [14]: `cat_data`

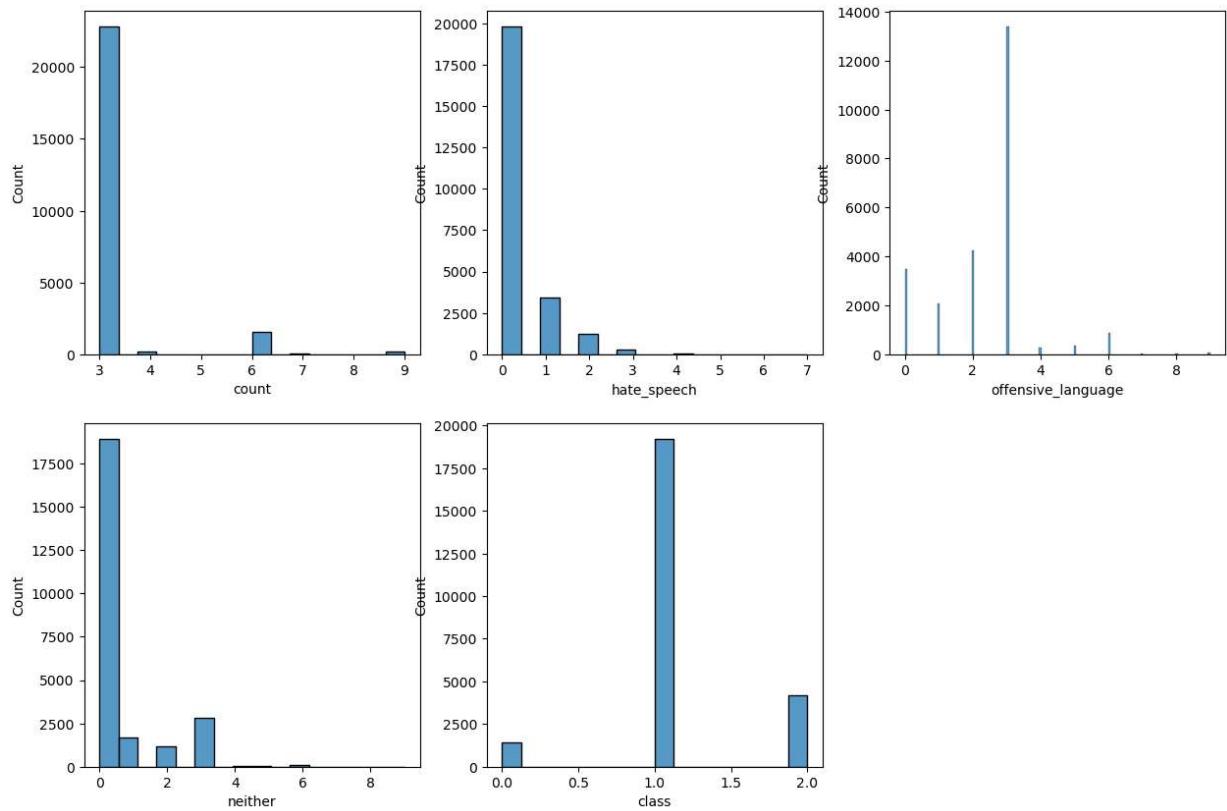
Out[14]: `['tweet', 'labels']`

## Visualization

In [15]: `import matplotlib.pyplot as plt  
import seaborn as sns`

```
In [16]: plt.figure(figsize=(15,10))
```

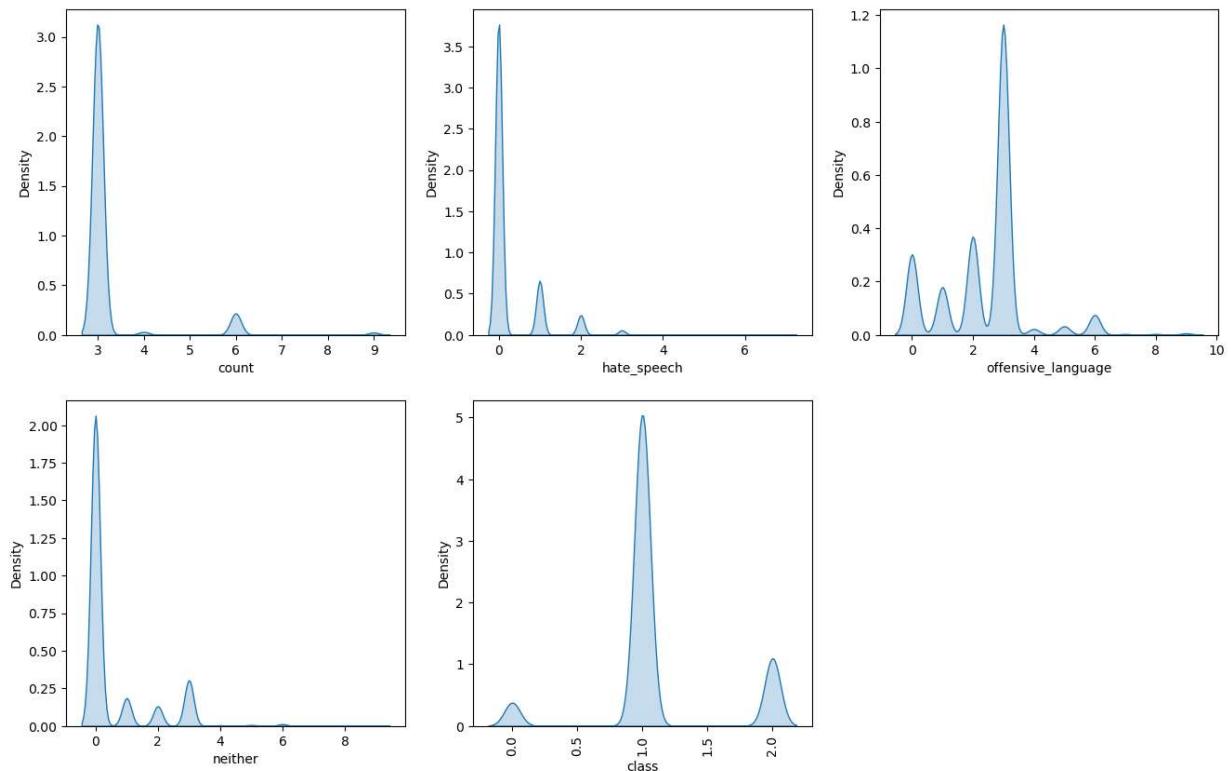
```
for i , col in enumerate(num_data):
    plt.subplot(2,3,i+1)
    sns.histplot(data=df,x=col)
plt.show()
```



## KDE Plot

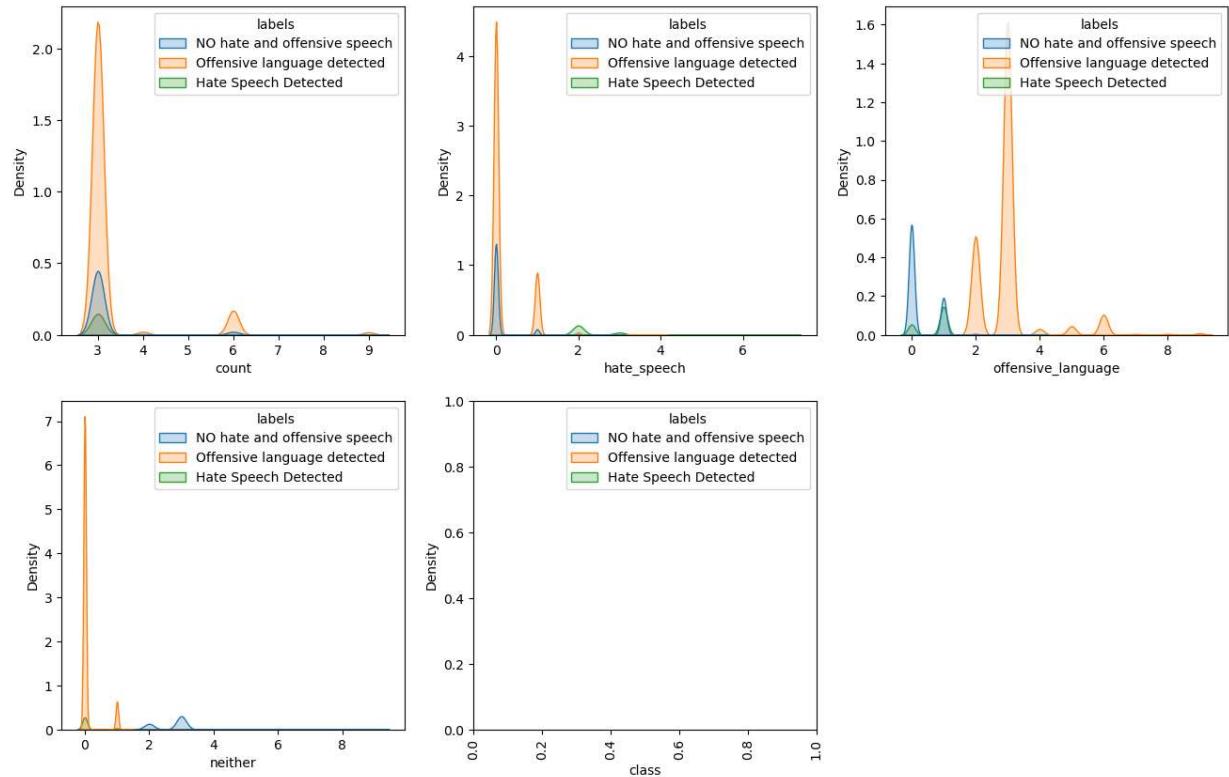
```
In [17]: plt.figure(figsize=(16,10))

for i , col in enumerate(num_data):
    plt.subplot(2,3,i+1)
    sns.kdeplot(data=df,x=col,fill=True)
plt.xticks(rotation = 90)
plt.show()
```



```
In [84]: plt.figure(figsize=(16,10))
```

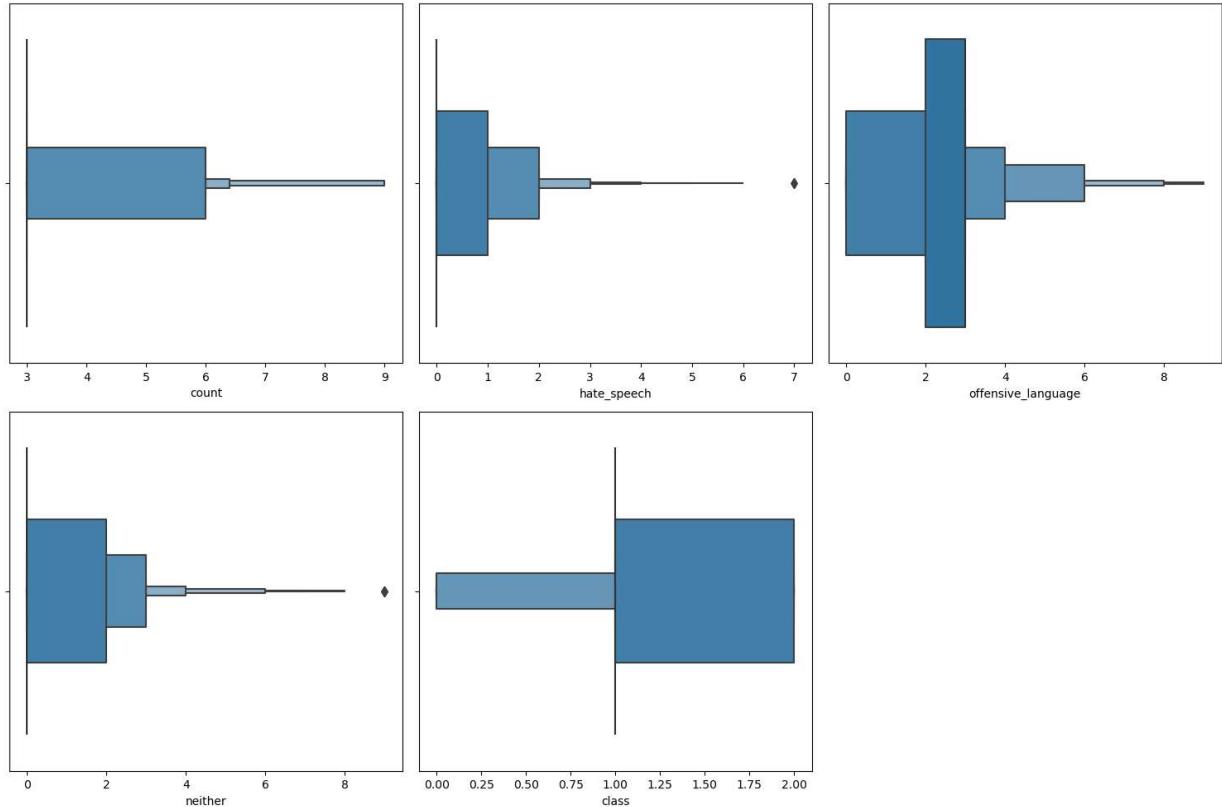
```
for i , col in enumerate(num_data):
    plt.subplot(2,3,i+1)
    sns.kdeplot(data=df,x=col,hue=df[ 'labels' ],fill=True)
plt.xticks(rotation = 90)
plt.show()
```



## Boxen Plot

```
In [18]: plt.figure(figsize=(15,10))
```

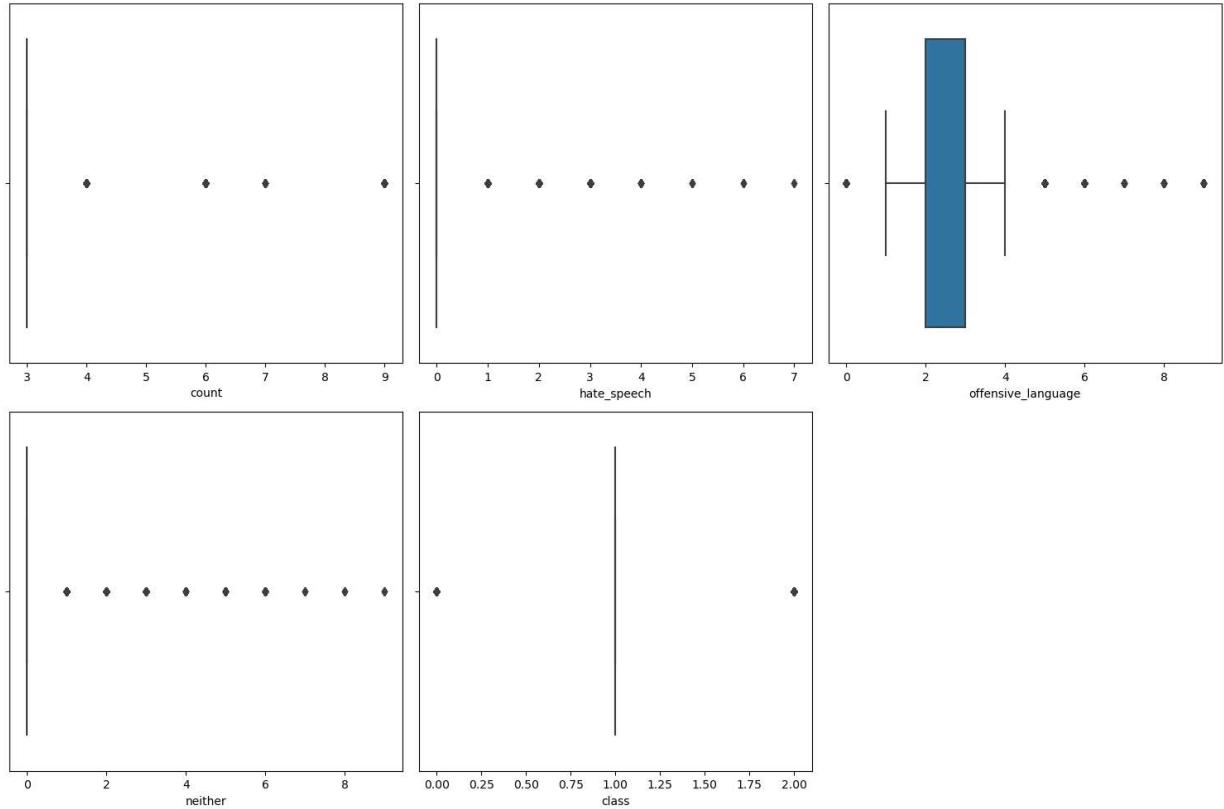
```
for i , col in enumerate(num_data):
    plt.subplot(2,3,i+1)
    sns.boxenplot(data=df,x=col)
plt.tight_layout()
plt.show()
```



## Box Plot

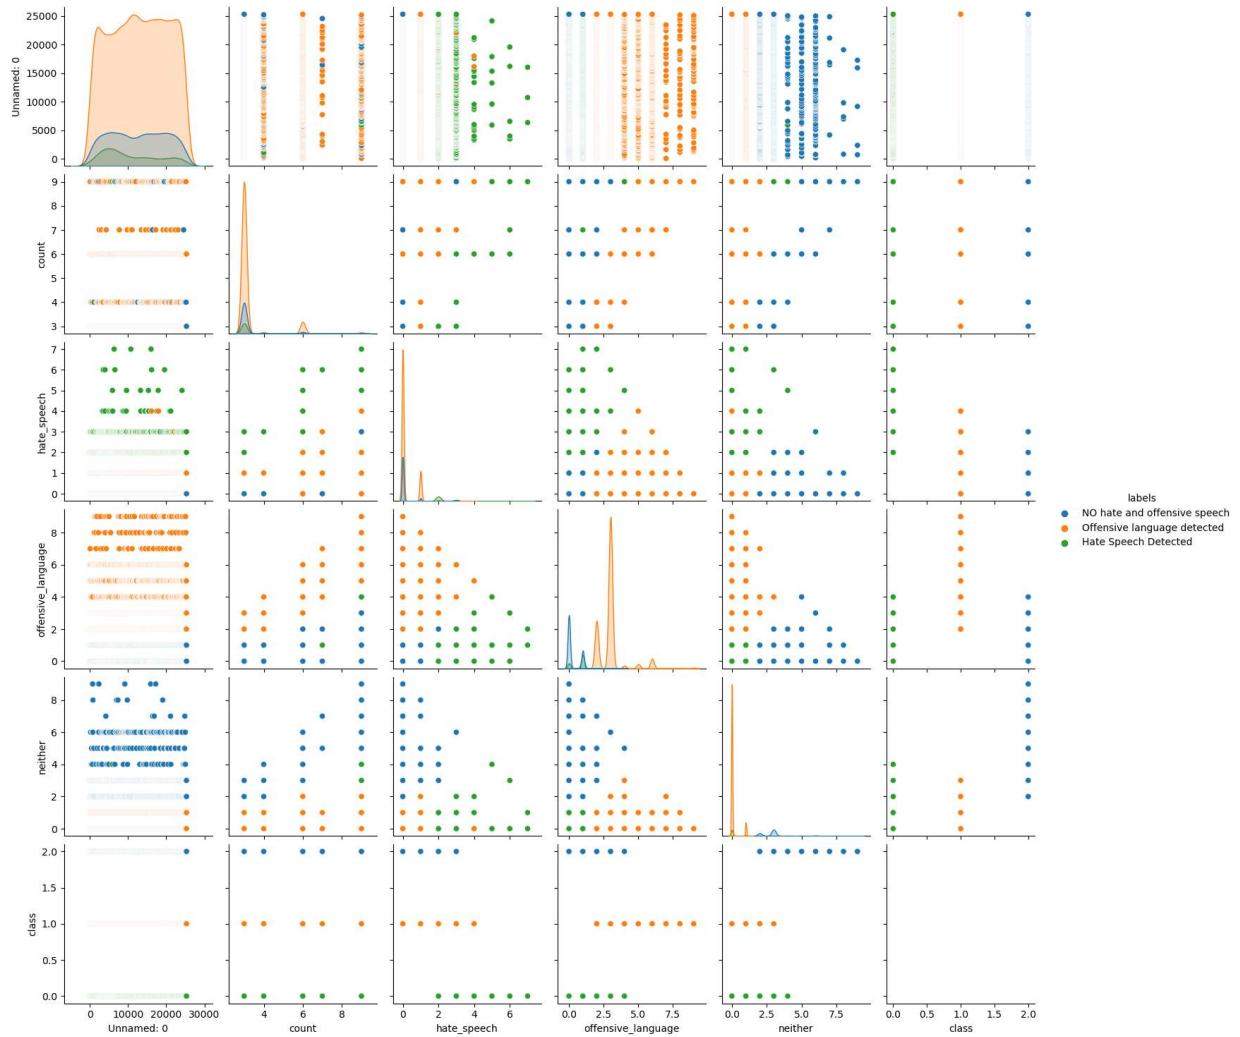
```
In [19]: plt.figure(figsize=(15,10))
```

```
for i , col in enumerate(num_data):
    plt.subplot(2,3,i+1)
    sns.boxplot(data=df,x=col)
plt.tight_layout()
plt.show()
```

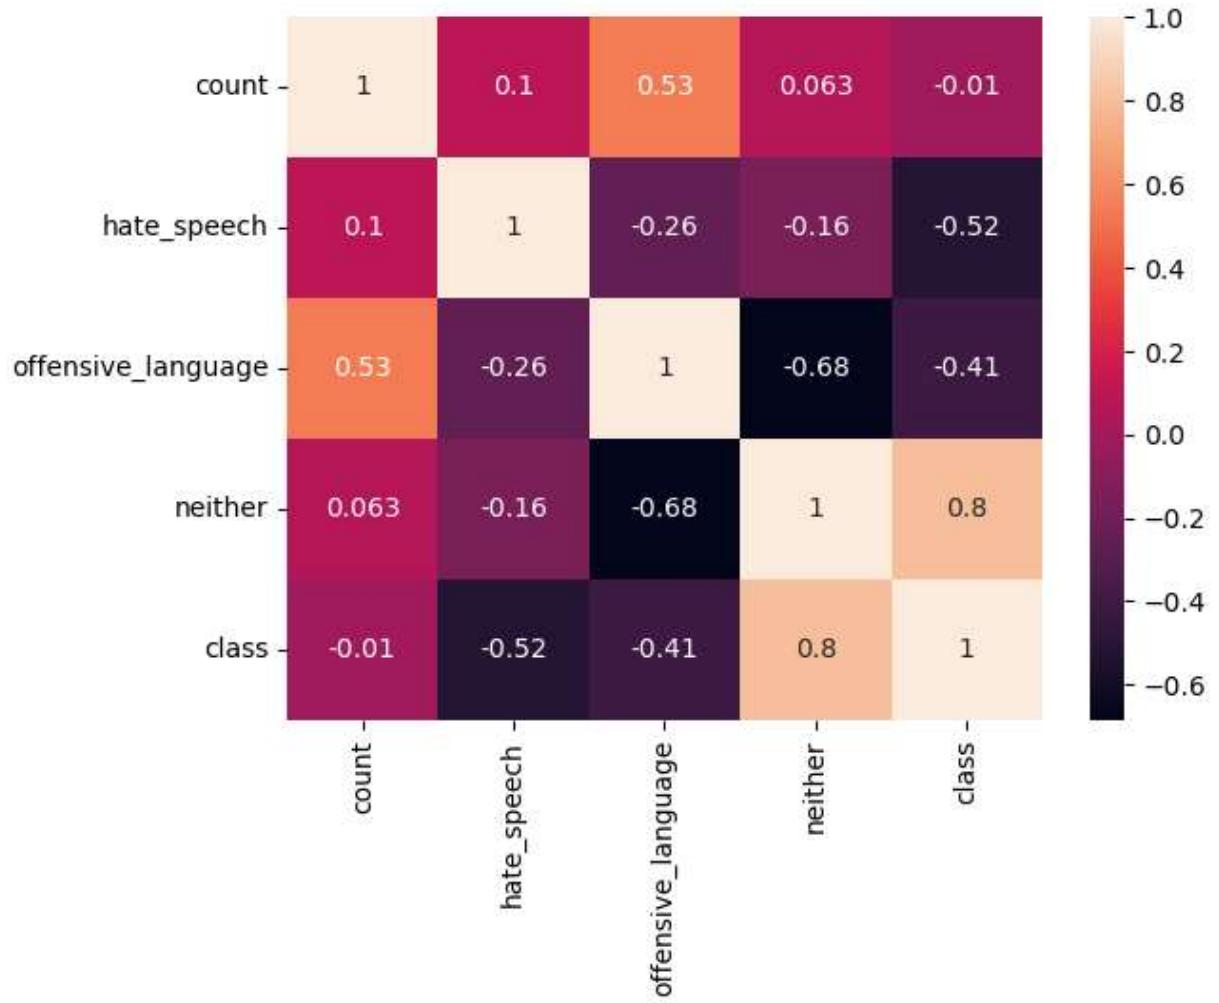


## Pair Plot

```
In [20]: sns.pairplot(df,hue='labels')
plt.show()
```

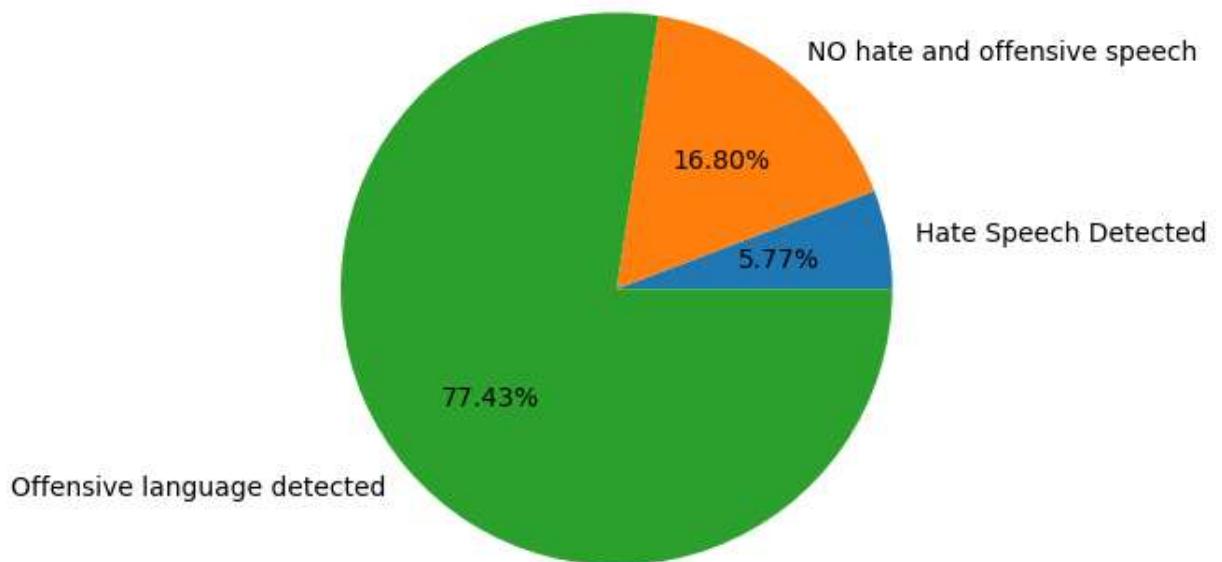


```
In [21]: sns.heatmap(df[num_data].corr(), annot=True)
plt.show()
```



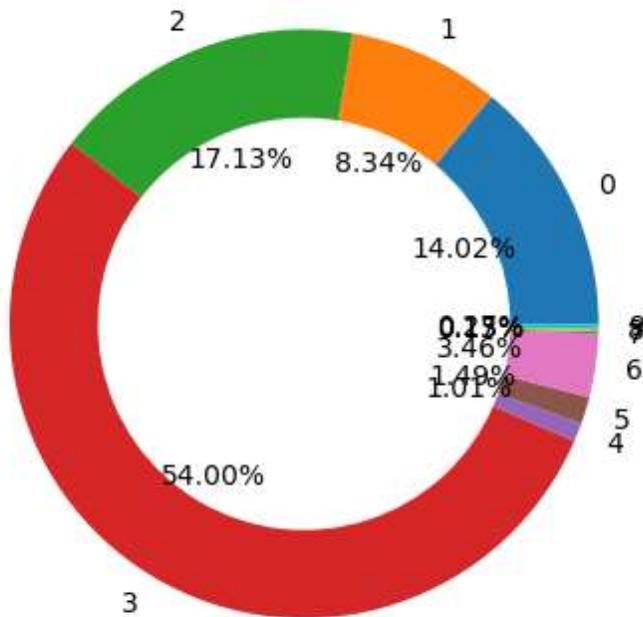
## Pie Chart

```
In [29]: h=df.groupby('labels')['labels'].count()  
plt.pie(h, labels=h.index, autopct='%.2f%%')  
plt.show()
```



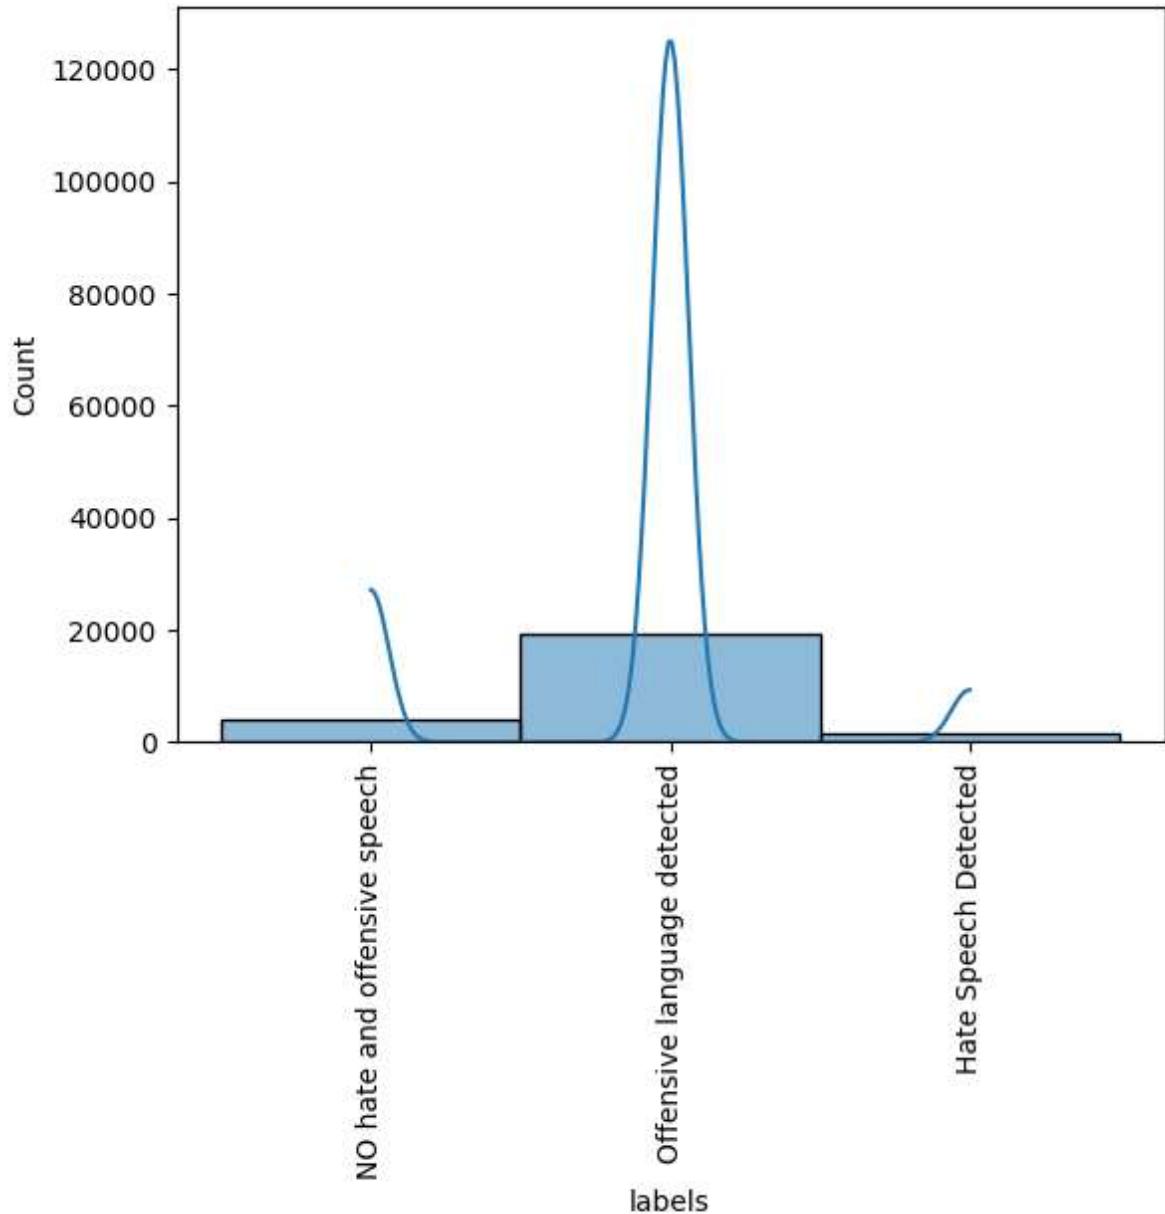
## Donut Chart

```
In [38]: i=df.groupby('offensive_language')['offensive_language'].count()  
plt.pie(i, labels=i.index, autopct='%.2f%%', wedgeprops=dict(width=0.3))  
plt.show()
```

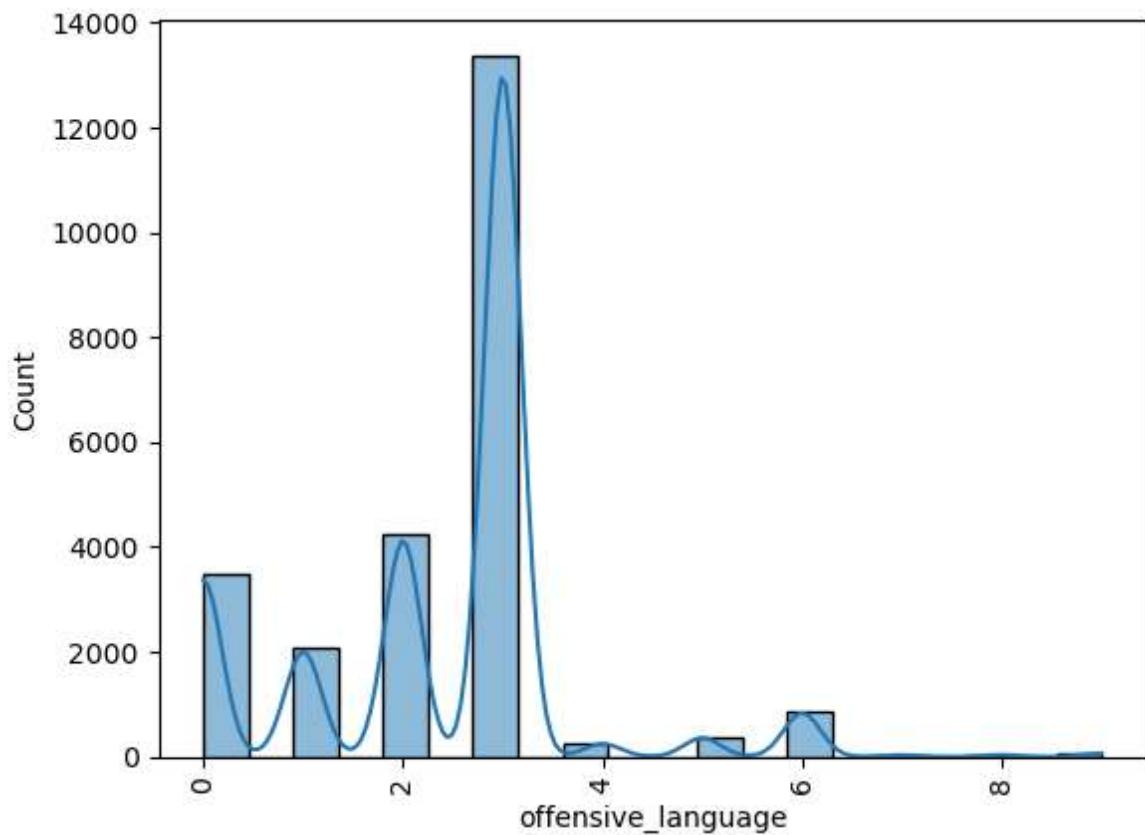


## Hist plot

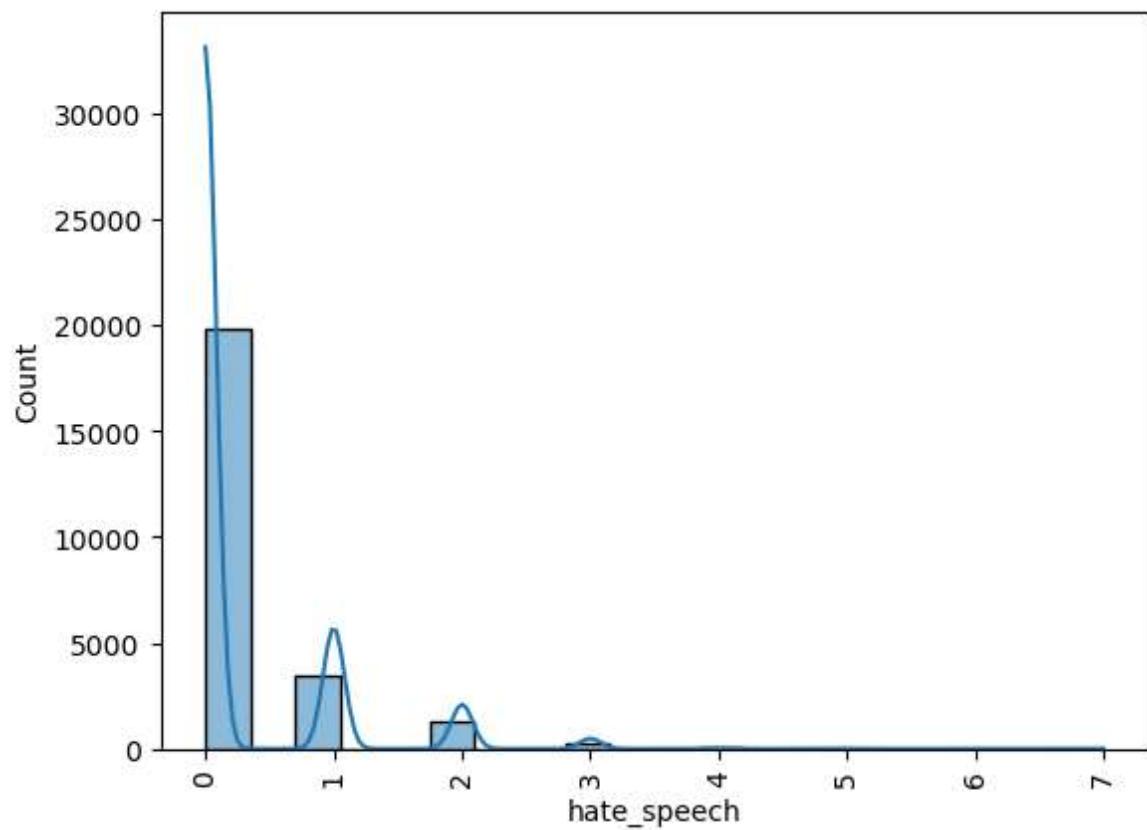
```
In [40]: sns.histplot(df['labels'], bins=20, kde=True)
plt.xticks(rotation=90)
plt.show()
```



```
In [42]: sns.histplot(df['offensive_language'], bins=20, kde=True)
plt.xticks(rotation=90)
plt.show()
```



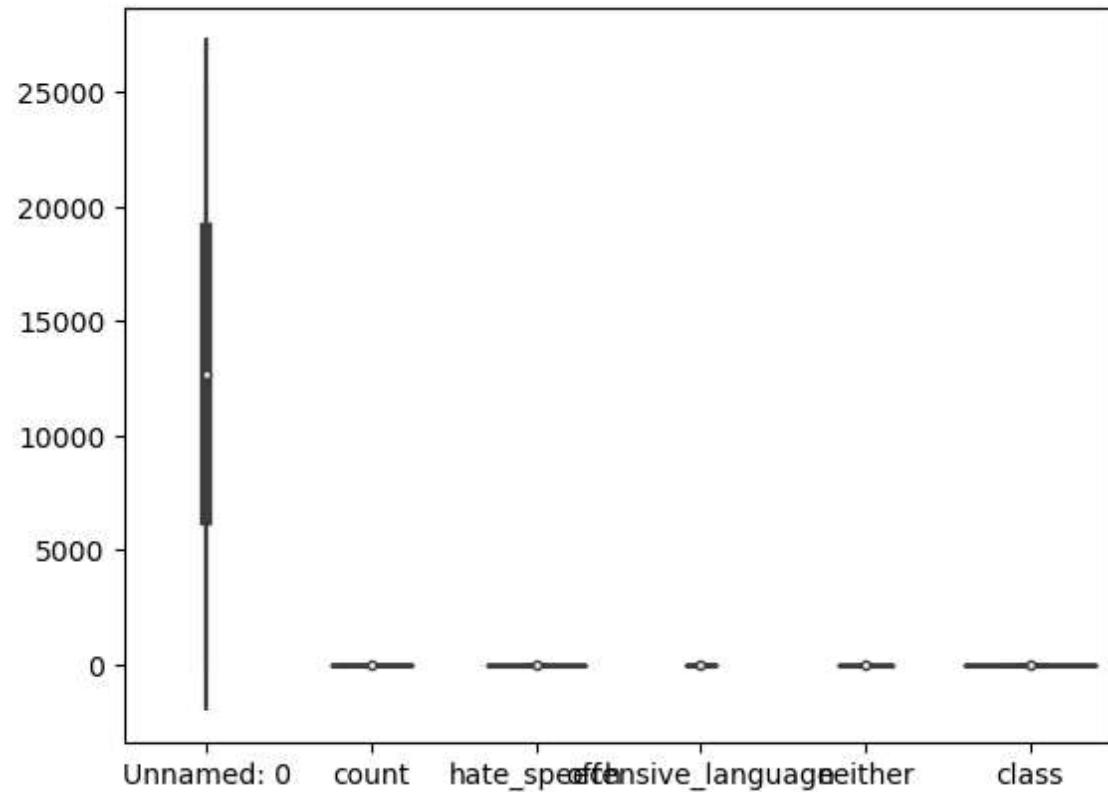
```
In [43]: sns.histplot(df['hate_speech'], bins=20, kde=True)
plt.xticks(rotation=90)
plt.show()
```



## Violinplot

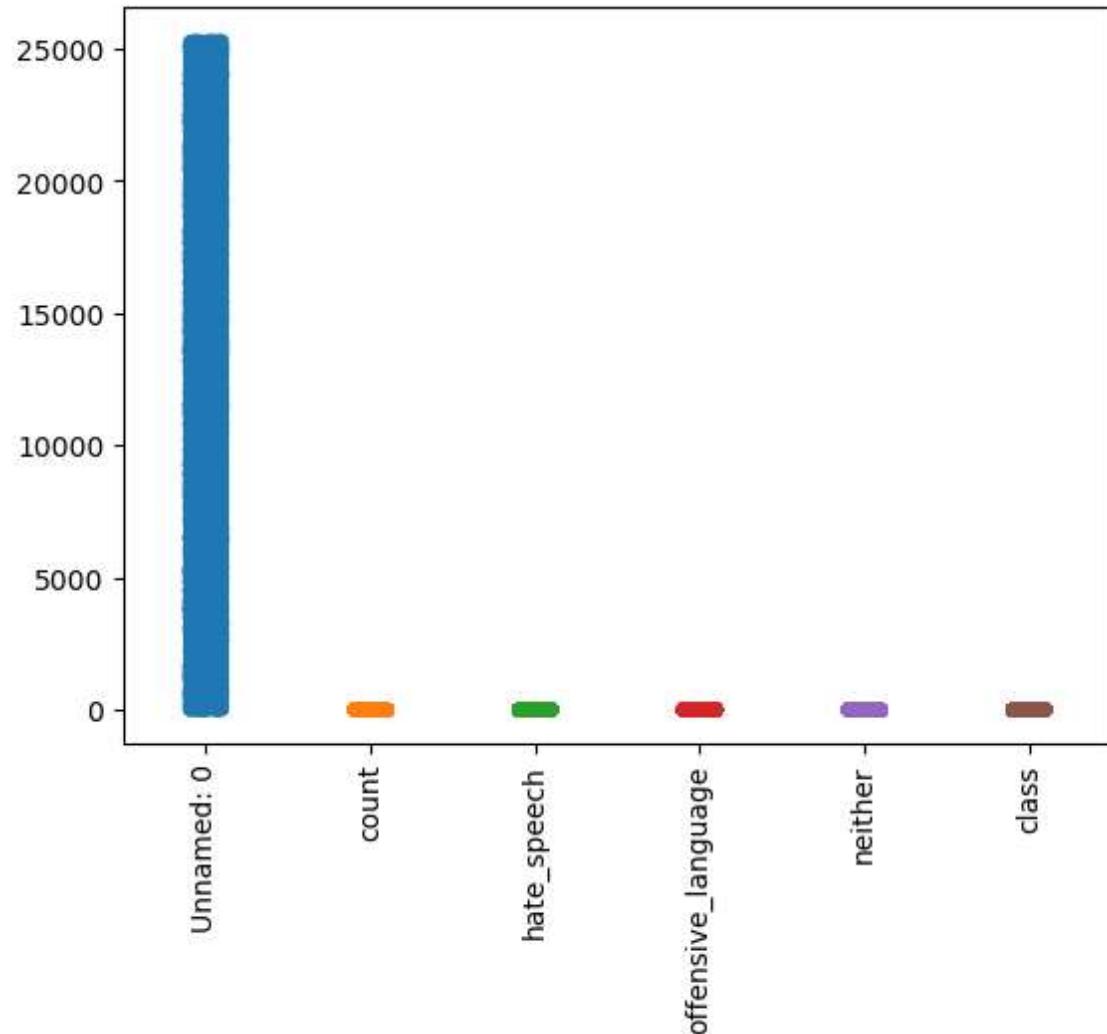
In [50]: `sns.violinplot(df)`

Out[50]: <Axes: >



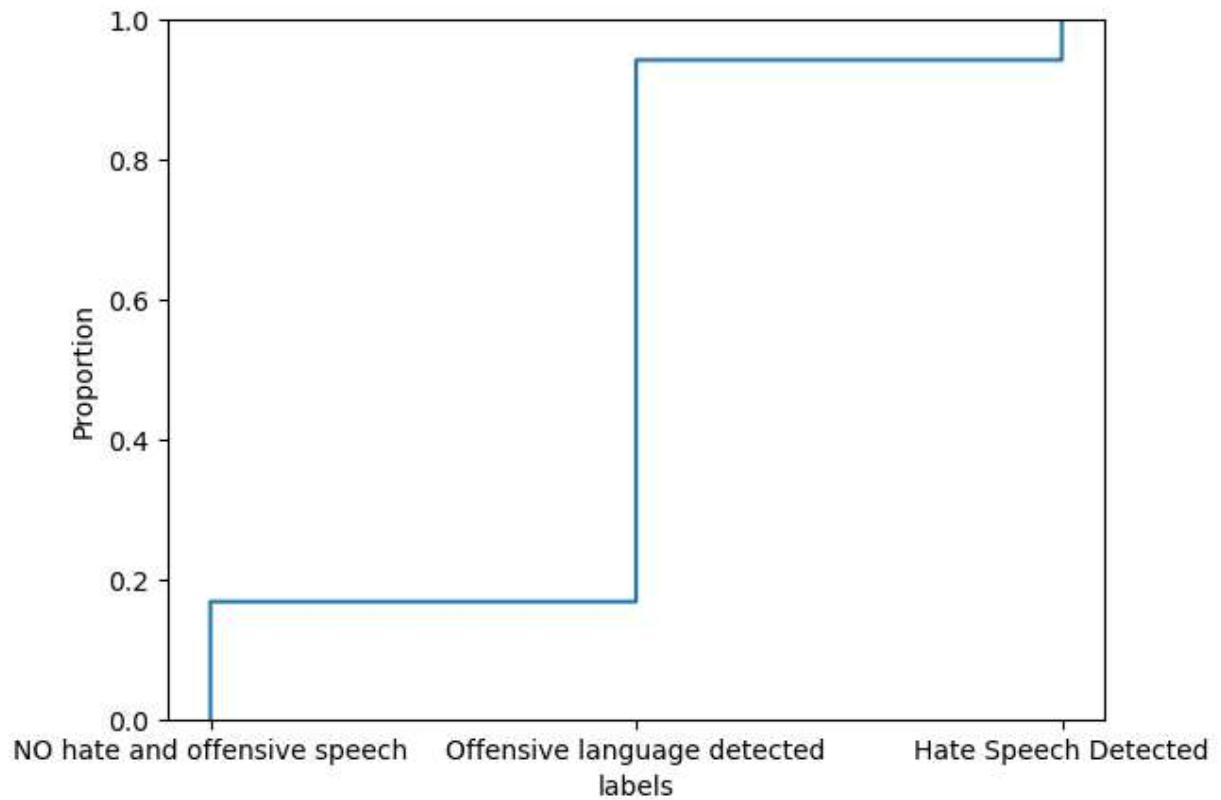
## Strip plot

```
In [55]: sns.stripplot(df)
plt.xticks(rotation=90)
plt.show()
```



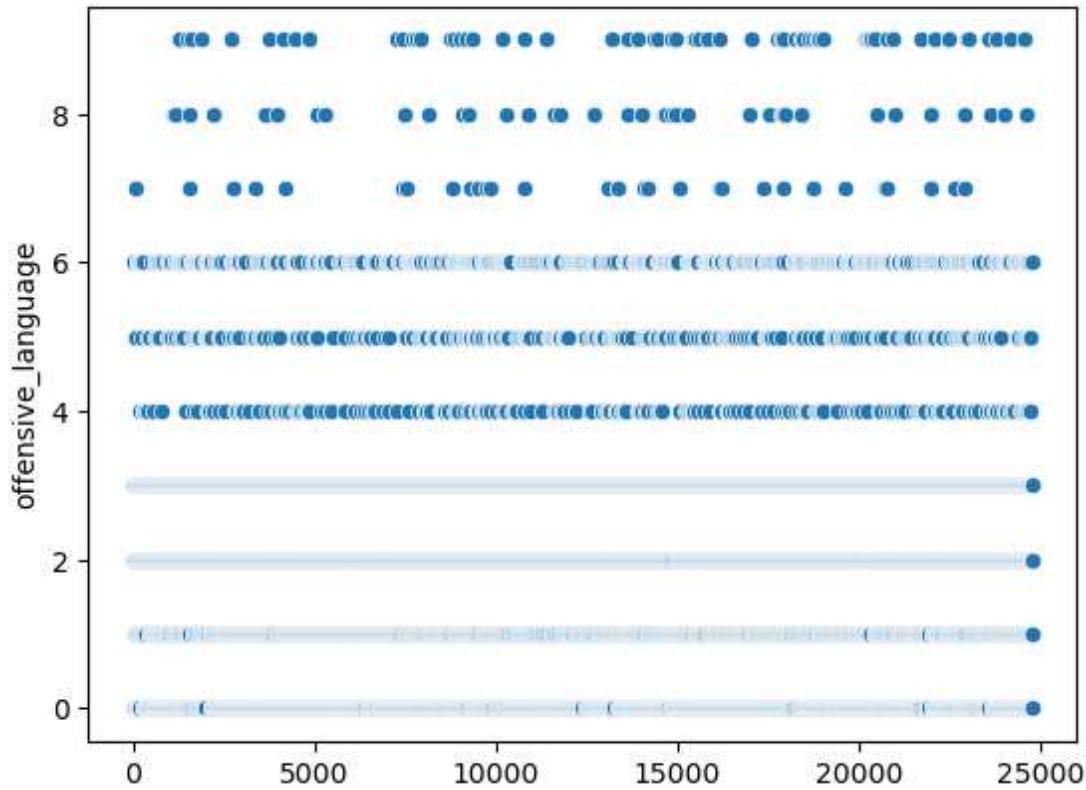
## ECDF Plot

```
In [70]: sns.ecdfplot(x=df['labels'], data=df)  
plt.show()
```



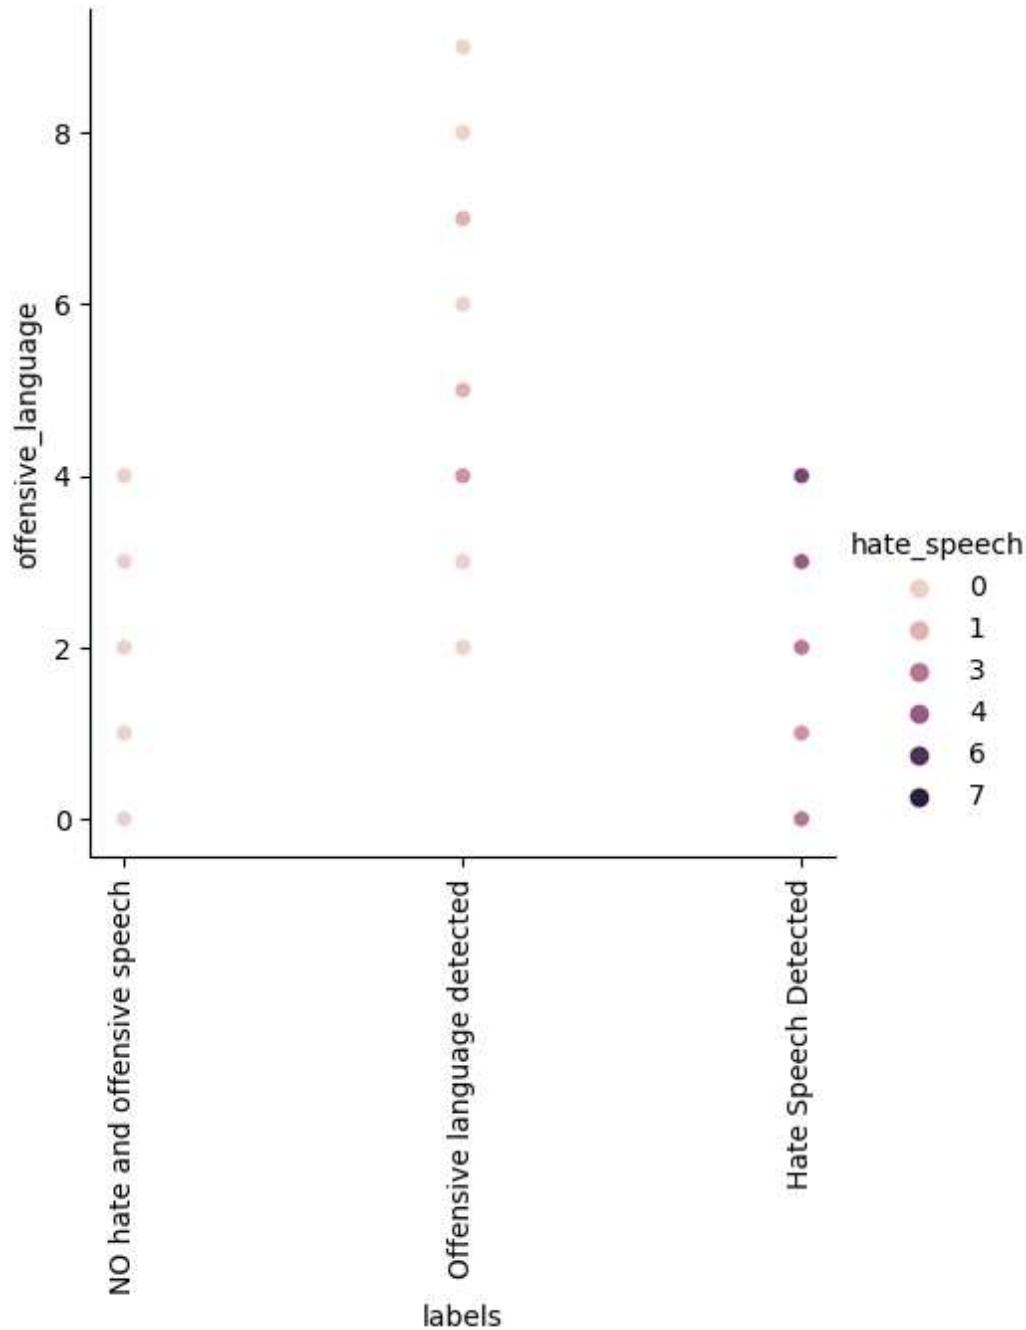
## Scatterplot

```
In [74]: sns.scatterplot(df['offensive_language'])
plt.show()
```



## Relplot

```
In [85]: sns.relplot(x='labels', y='offensive_language', hue = 'hate_speech', data=df)
plt.xticks(rotation=90)
plt.show()
```



```
In [86]: df=df[['tweet','labels']]
```

In [87]: df

Out[87]:

		tweet	labels
0	!!! RT @mayasolovel: As a woman you shouldn't...	NO hate and offensive speech	
1	!!!! RT @mleew17: boy dats cold...tyga dwn ba...	Offensive language detected	
2	!!!!!! RT @UrKindOfBrand Dawg!!!! RT @80sbaby...	Offensive language detected	
3	!!!!!!! RT @C_G_Anderson: @viva_based she lo...	Offensive language detected	
4	!!!!!!!!!! RT @ShenikaRoberts: The shit you...	Offensive language detected	
...	...	...	...
24778	you's a muthaf***in lie &#8220;@LifeAsKing: @2...	Offensive language detected	
24779	you've gone and broke the wrong heart baby, an...	NO hate and offensive speech	
24780	young buck wanna eat!!.. dat nigguh like I ain...	Offensive language detected	
24781	youu got wild bitches tellin you lies	Offensive language detected	
24782	~~Ruffled   Ntac Eileen Dahlia - Beautiful col...	NO hate and offensive speech	

24783 rows × 2 columns

In [88]:

```
def clean(text):
    text = str(text).lower()
    text = re.sub('\[.*?\]', '', text)
    text = re.sub('https?://\S+|www\.\S+', '', text)
    text = re.sub('<.*?>+', '', text)
    text = re.sub('[%s]' % re.escape(string.punctuation), ' ', text)
    text = re.sub('\n', ' ', text)
    text = re.sub('\w*\d\w*', ' ', text)
    text = [word for word in text.split(' ') if word not in stopword]
    text = " ".join(text)
    text = [stemmer.stem(word) for word in text.split(' ')]
    text = " ".join(text)
    return text

df["tweet"] = df["tweet"].apply(clean)
print(df.head())
```

	tweet \
0	rt mayasolov woman shouldnt complain clean ho...
1	rt boy dat coldtyga dwn bad cuffin dat hoe ...
2	rt urkindofbrand dawg rt ever fuck bitch sta...
3	rt cganderson vivabas look like tranni
4	rt shenikarobert shit hear might true might f...

	labels
0	NO hate and offensive speech
1	Offensive language detected
2	Offensive language detected
3	Offensive language detected
4	Offensive language detected

In [89]: `df.isnull().sum()`

Out[89]:  
 tweet 0  
 labels 0  
 dtype: int64

In [90]: `df.head()`

Out[90]:

	tweet	labels
0	rt mayasolov woman shouldnt complain clean ho...	NO hate and offensive speech
1	rt boy dat coldtyga dwn bad cuffin dat hoe ...	Offensive language detected
2	rt urkindofbrand dawg rt ever fuck bitch sta...	Offensive language detected
3	rt cganderson vivabas look like tranni	Offensive language detected
4	rt shenikarobert shit hear might true might f...	Offensive language detected

## Decision Tree

In [91]: `x=np.array(df["tweet"])`  
`y=np.array(df["labels"])`

In [92]: `cv=CountVectorizer()`

In [93]: `x=cv.fit_transform(x)`

In [94]: `x`

Out[94]: <24783x25693 sparse matrix of type '<class 'numpy.int64'>'  
 with 197862 stored elements in Compressed Sparse Row format>

In [95]: `y`

Out[95]: `array(['NO hate and offensive speech', 'Offensive language detected', 'Offensive language detected', ..., 'Offensive language detected', 'Offensive language detected', 'NO hate and offensive speech'], dtype=object)`

In [96]: `x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.33,random_state=42)`

In [97]: `clf=DecisionTreeClassifier()`

```
In [98]: clf.fit(x_train,y_train)
```

```
Out[98]:
```

```
  ▾ DecisionTreeClassifier
    DecisionTreeClassifier()
```

```
In [99]: y_pred=clf.predict(x_test)
```

```
In [100]: from sklearn.metrics import accuracy_score, classification_report
```

```
acc = accuracy_score(y_pred,y_test)
print(acc)
```

```
0.8784692505196234
```

```
In [101]: print('Classification Report')
print(classification_report(y_pred,y_test))
```

	precision	recall	f1-score	support
Hate Speech Detected	0.31	0.37	0.34	399
NO hate and offensive speech	0.82	0.82	0.82	1388
Offensive language detected	0.93	0.92	0.93	6392
accuracy			0.88	8179
macro avg	0.69	0.70	0.70	8179
weighted avg	0.88	0.88	0.88	8179

```
In [102]: test_data = "I will kill you"
df1=cv.transform([test_data]).toarray()
print(clf.predict(df1))
```

```
['Hate Speech Detected']
```

```
In [103]: test_data2 = "You are bad i don't like you"
df2=cv.transform([test_data2]).toarray()
print(clf.predict(df2))
```

```
['Offensive language detected']
```

```
In [104]: test_data3 = "I will not kill you"
df3=cv.transform([test_data3]).toarray()
print(clf.predict(df3))
```

```
['Hate Speech Detected']
```

## KNN

```
In [360]: from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix
```

```
In [361]: x_knn=np.array(df["tweet"])
y_knn=np.array(df["labels"])
```

```
In [362]: cv_knn=CountVectorizer()
```

```
In [363]: x_knn= cv.fit_transform(x_knn)
```

```
In [364]: X_train_knn, X_test_knn, Y_train_knn, Y_test_knn = train_test_split(x_knn, y_knn,
```

```
In [365]: knn = KNeighborsClassifier(n_neighbors=10)

#Training the model
knn.fit(X_train_knn,Y_train_knn)
```

Out[365]:

```
▼ KNeighborsClassifier
KNeighborsClassifier(n_neighbors=10)
```

```
In [366]: knn.score(X_test_knn,Y_test_knn)
```

Out[366]: 0.8482953399233407

```
In [367]: #Making prediction on test set
Y_pred_knn= knn.predict(X_test_knn)
```

```
In [368]: #Accuracy of KNN model
acc_knn = accuracy_score(Y_test_knn, Y_pred_knn)
print("Accuracy:", acc_knn)
```

Accuracy: 0.8482953399233407

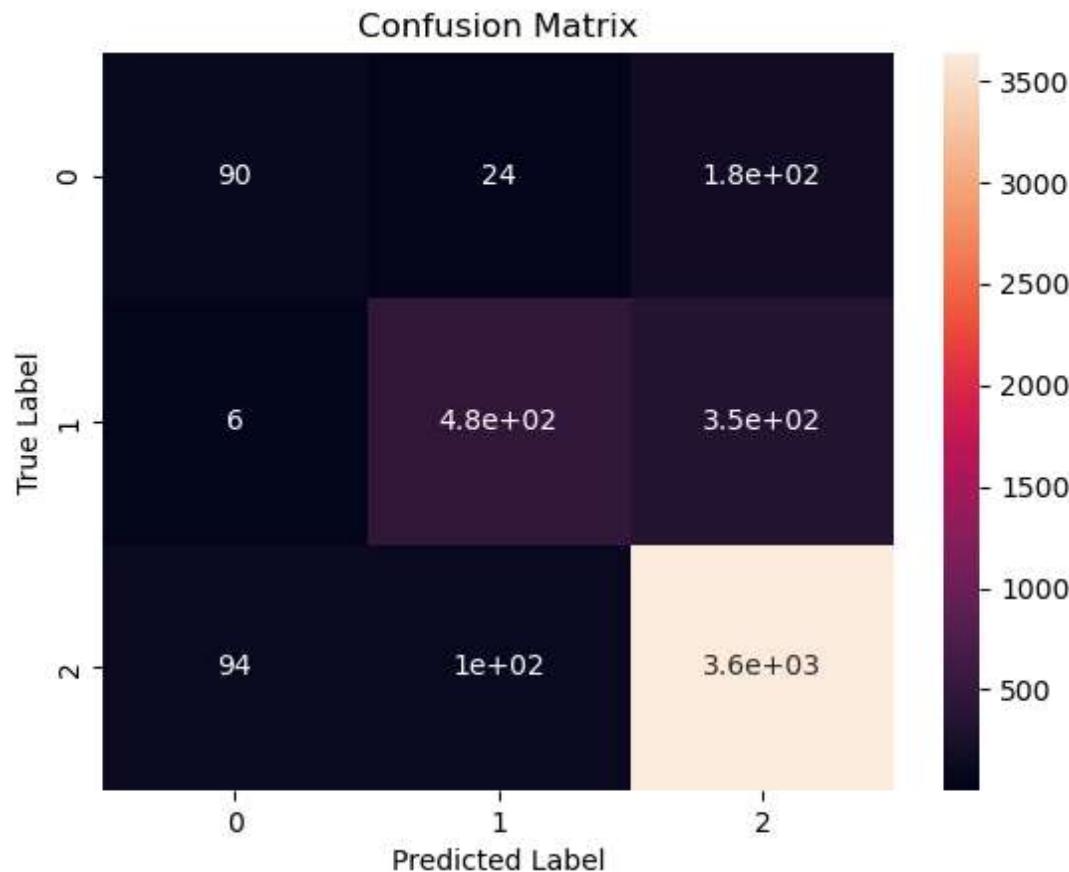
```
In [369]: #Classification report of KNN model
print(classification_report(Y_test_knn, Y_pred_knn))
```

	precision	recall	f1-score	support
Hate Speech Detected	0.47	0.31	0.38	290
NO hate and offensive speech	0.79	0.57	0.67	835
Offensive language detected	0.87	0.95	0.91	3832
accuracy			0.85	4957
macro avg	0.71	0.61	0.65	4957
weighted avg	0.84	0.85	0.84	4957

```
In [370]: #Confusion Matrix of KNN model  
cm=confusion_matrix(Y_test_knn, Y_pred_knn)  
print('Confusion matrix : ')  
print(cm)  
print('True Positives(TP) = ', cm[0,0])  
print('True Negatives(TN) = ', cm[1,1])  
print('False Positives(FP) = ', cm[0,1])  
print('False Negatives(FN) = ', cm[1,0])
```

```
Confusion matrix :  
[[ 90  24 176]  
 [  6 479 350]  
 [ 94 102 3636]]  
True Positives(TP) =  90  
True Negatives(TN) =  479  
False Positives(FP) =  24  
False Negatives(FN) =  6
```

```
In [371]: sns.heatmap(cm, annot=True)  
plt.title('Confusion Matrix')  
plt.xlabel('Predicted Label')  
plt.ylabel('True Label')  
plt.show()
```



## KNN Tuning

```
In [117]: from sklearn.model_selection import GridSearchCV
```

```
In [118]: #Grid Search
```

```
param_grid = {
    'n_neighbors': [3, 5, 7, 9], # Test different values of k
    'weights': ['uniform', 'distance'], # Weighting type
    'p': [1, 2], # Distance metric (1 for Manhattan, 2 for Euclidean)
}
```

```
In [119]: knn = KNeighborsClassifier()
grid_search = GridSearchCV(knn, param_grid, cv=5, scoring='accuracy')
grid_search.fit(X_train_knn, Y_train_knn)
```

Out[119]:

```
▶ GridSearchCV
  ▶ estimator: KNeighborsClassifier
    ▶ KNeighborsClassifier
```

```
In [120]: best_param = grid_search.best_params_
best_knn = KNeighborsClassifier(n_neighbors = best_param['n_neighbors'], weights = best_param['weights'],
best_knn.fit(X_train_knn, Y_train_knn)
Y_pred_knn = best_knn.predict(X_test_knn)
```

```
In [121]: print("Best Hyperparameter : ", best_param)
```

```
Best Hyperparameter : {'n_neighbors': 5, 'p': 1, 'weights': 'distance'}
```

```
In [122]: acc_knn = accuracy_score(Y_test_knn, Y_pred_knn)
acc_knn
```

Out[122]: 0.8539439176921525

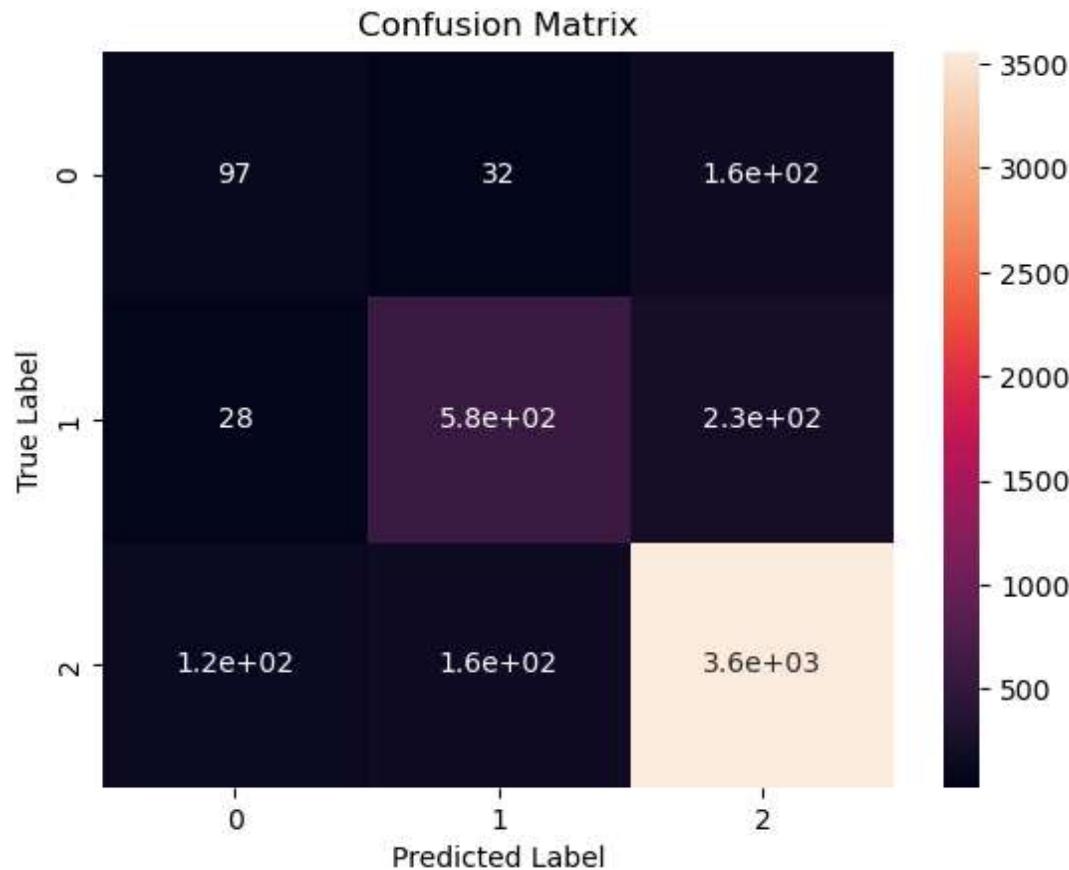
```
In [123]: print (classification_report(Y_test_knn, Y_pred_knn))
```

	precision	recall	f1-score	support
Hate Speech Detected	0.40	0.33	0.37	290
NO hate and offensive speech	0.75	0.69	0.72	835
Offensive language detected	0.90	0.93	0.91	3832
accuracy			0.85	4957
macro avg	0.69	0.65	0.67	4957
weighted avg	0.85	0.85	0.85	4957

```
In [124]: cm=confusion_matrix(Y_test_knn, Y_pred_knn)
print('Confusion matrix : ')
print(cm)
print('True Positives(TP) = ', cm[0,0])
print('True Negatives(TN) = ', cm[1,1])
print('False Positives(FP) = ', cm[0,1])
print('False Negatives(FN) = ', cm[1,0])
```

```
Confusion matrix :
[[ 97  32 161]
 [ 28 579 228]
 [ 116 159 3557]]
True Positives(TP) = 97
True Negatives(TN) = 579
False Positives(FP) = 32
False Negatives(FN) = 28
```

```
In [125]: sns.heatmap(cm, annot=True)
plt.title('Confusion Matrix')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```



In [126]: #Euclidean method

```
knn_classifier = KNeighborsClassifier(n_neighbors=10, metric='euclidean')
```

In [127]: knn\_classifier.fit(X\_train\_knn, Y\_train\_knn)

Out[127]:

```
KNeighborsClassifier
```

```
KNeighborsClassifier(metric='euclidean', n_neighbors=10)
```

In [128]: Y\_pred\_knn\_eu = knn\_classifier.predict(X\_test\_knn)

In [129]: accuracy = accuracy\_score(Y\_test\_knn, Y\_pred\_knn\_eu) # For classification  
print('accuracy:', accuracy)

```
accuracy: 0.8482953399233407
```

In [130]: print(classification\_report(Y\_test\_knn, Y\_pred\_knn\_eu))

	precision	recall	f1-score	support
Hate Speech Detected	0.47	0.31	0.38	290
NO hate and offensive speech	0.79	0.57	0.67	835
Offensive language detected	0.87	0.95	0.91	3832
accuracy			0.85	4957
macro avg	0.71	0.61	0.65	4957
weighted avg	0.84	0.85	0.84	4957

In [131]: cm=confusion\_matrix(Y\_test\_knn, Y\_pred\_knn\_eu)

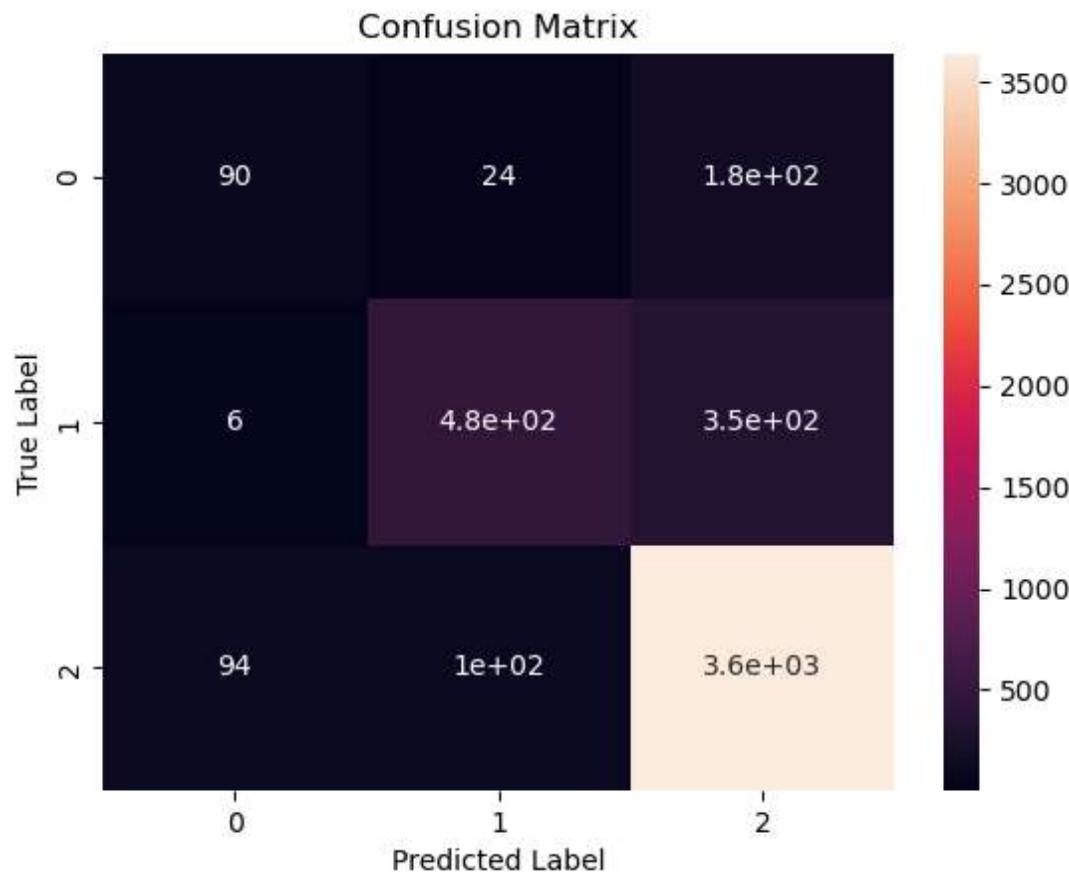
```
Confusion matrix :  
[[ 90  24 176]  
 [  6 479 350]  
 [ 94 102 3636]]  
True Positives(TP) =  90  
True Negatives(TN) = 479  
False Positives(FP) =  24  
False Negatives(FN) =   6
```

Confusion matrix :

```
[[ 90  24 176]  
 [  6 479 350]  
 [ 94 102 3636]]
```

```
True Positives(TP) = 90  
True Negatives(TN) = 479  
False Positives(FP) = 24  
False Negatives(FN) = 6
```

```
In [132]: sns.heatmap(cm, annot=True)
plt.title('Confusion Matrix')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```



```
In [133]: #Manhattan distance method

knn_classifier = KNeighborsClassifier(n_neighbors=10, metric='manhattan')
```

```
In [134]: knn_classifier.fit(X_train_knn, Y_train_knn)
```

Out[134]:

```
▼          KNeighborsClassifier
KNeighborsClassifier(metric='manhattan', n_neighbors=10)
```

```
In [135]: Y_pred_knn_mh = knn_classifier.predict(X_test_knn)
```

```
In [136]: accuracy_knn_mh = accuracy_score(Y_test_knn, Y_pred_knn_mh)
print('accuracy:', accuracy_knn_mh)
```

accuracy: 0.8452693161186201

```
In [137]: print(classification_report(Y_test_knn, Y_pred_knn))
```

	precision	recall	f1-score	support
Hate Speech Detected	0.40	0.33	0.37	290
NO hate and offensive speech	0.75	0.69	0.72	835
Offensive language detected	0.90	0.93	0.91	3832
accuracy			0.85	4957
macro avg	0.69	0.65	0.67	4957
weighted avg	0.85	0.85	0.85	4957

```
In [138]: cm=confusion_matrix(Y_test_knn, Y_pred_knn_mh)
```

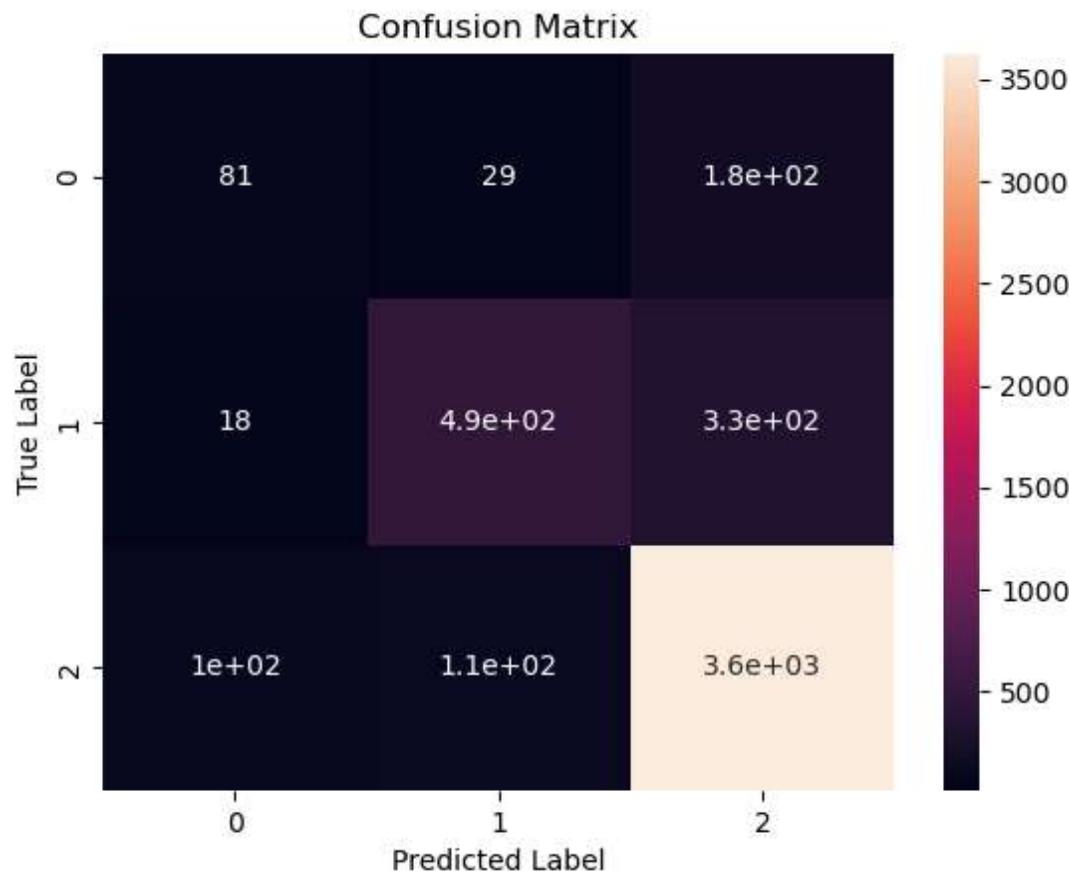
```
print('Confusion matrix : ')
print(cm)
print('True Positives(TP) = ', cm[0,0])
print('True Negatives(TN) = ', cm[1,1])
print('False Positives(FP) = ', cm[0,1])
print('False Negatives(FN) = ', cm[1,0])
```

```
Confusion matrix :
```

```
[[ 81  29 180]
 [ 18 486 331]
 [102 107 3623]]
```

```
True Positives(TP) = 81
True Negatives(TN) = 486
False Positives(FP) = 29
False Negatives(FN) = 18
```

```
In [139]: sns.heatmap(cm, annot=True)
plt.title('Confusion Matrix')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```



```
In [140]: #Minkowsky distance method
```

```
knn_classifier = KNeighborsClassifier(n_neighbors=10, metric='minkowski', p=2)
```

```
In [141]: knn_classifier.fit(X_train_knn, Y_train_knn)
```

```
Out[141]:
```

```
▼ KNeighborsClassifier
KNeighborsClassifier(n_neighbors=10)
```

```
In [142]: Y_pred_knn_mk = knn_classifier.predict(X_test_knn)
```

```
In [143]: accuracy_knn_mk = accuracy_score(Y_test_knn, Y_pred_knn_mk) # For classification
print('accuracy:', accuracy_knn_mk)
```

```
accuracy: 0.8482953399233407
```

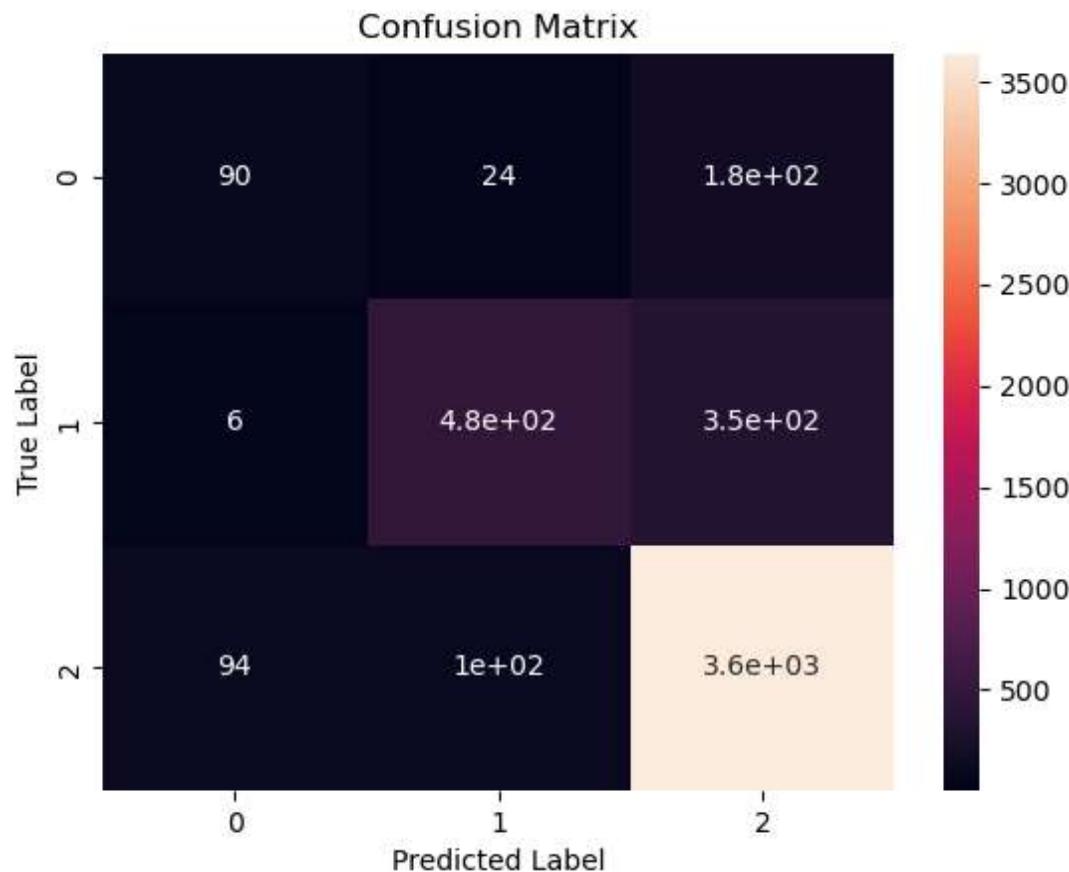
```
In [144]: print(classification_report(Y_test_knn, Y_pred_knn_mk))
```

	precision	recall	f1-score	support
Hate Speech Detected	0.47	0.31	0.38	290
NO hate and offensive speech	0.79	0.57	0.67	835
Offensive language detected	0.87	0.95	0.91	3832
accuracy			0.85	4957
macro avg	0.71	0.61	0.65	4957
weighted avg	0.84	0.85	0.84	4957

```
In [145]: cm=confusion_matrix(Y_test_knn, Y_pred_knn_mk)
print('Confusion matrix : ')
print(cm)
print('True Positives(TP) = ', cm[0,0])
print('True Negatives(TN) = ', cm[1,1])
print('False Positives(FP) = ', cm[0,1])
print('False Negatives(FN) = ', cm[1,0])
```

```
Confusion matrix :
[[ 90  24 176]
 [  6 479 350]
 [ 94 102 3636]]
True Positives(TP) =  90
True Negatives(TN) =  479
False Positives(FP) =  24
False Negatives(FN) =  6
```

```
In [146]: sns.heatmap(cm, annot=True)
plt.title('Confusion Matrix')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```



## SVM

```
In [147]: from sklearn.svm import SVC
from sklearn.multiclass import OneVsRestClassifier, OneVsOneClassifier
```

```
In [148]: x_svm=np.array(df["tweet"])
y_svm=np.array(df["labels"])
```

```
In [149]: cv_svm=CountVectorizer()
x_svm= cv.fit_transform(x_svm)
```

```
In [150]: X_train_svm, X_test_svm, Y_train_svm, Y_test_svm = train_test_split(x_svm, y_svm,
```

```
In [151]: #Running with default hyperparameter
svc = SVC()
```

In [152]: `svc.fit(X_train_svm, Y_train_svm)`

Out[152]:

```
    ▾ SVC
      SVC()
```

In [153]: `Y_pred_svm = svc.predict(X_test_svm)`

In [154]: *#Accuracy of SVM model*

```
acc_svm = accuracy_score(Y_test_svm, Y_pred_svm)
print("Accuracy Score with default parameter : {:.2f}%". format(acc_svm*100))
```

Accuracy Score with default parameter : 89.29%

In [155]: *#Classification report of SVM model*

```
print(classification_report(Y_test_svm, Y_pred_svm))
```

	precision	recall	f1-score	support
Hate Speech Detected	0.43	0.07	0.12	290
NO hate and offensive speech	0.80	0.90	0.85	835
Offensive language detected	0.92	0.95	0.94	3832
accuracy			0.89	4957
macro avg	0.71	0.64	0.63	4957
weighted avg	0.87	0.89	0.87	4957

In [156]: *#Confusion Matrix of SVM model*

```
cm = confusion_matrix(Y_test_svm, Y_pred_svm)
print('Confusion matrix : ')
print(cm)
print('True Positives(TP) = ', cm[0,0])
print('True Negatives(TN) = ', cm[1,1])
print('False Positives(FP) = ', cm[0,1])
print('False Negatives(FN) = ', cm[1,0])
```

Confusion matrix :

```
[[ 20   38  232]
 [  1  755   79]
 [ 26  155 3651]]
```

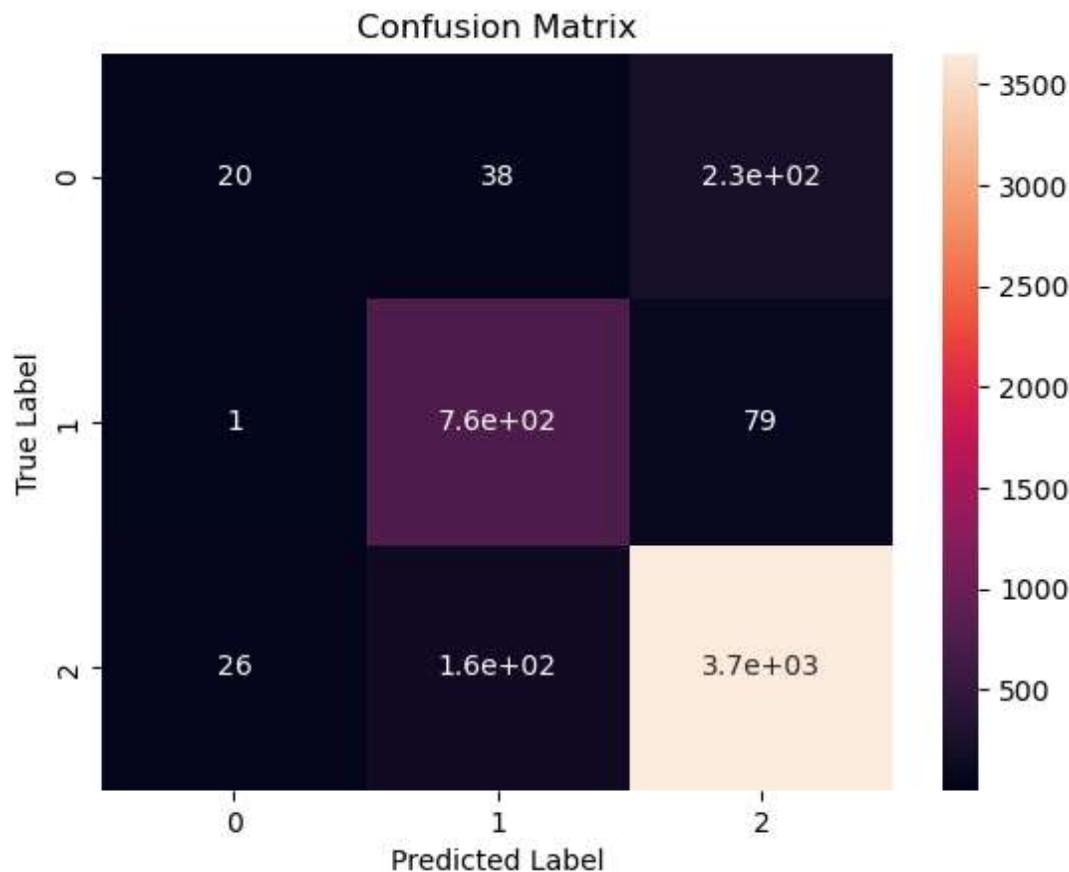
True Positives(TP) = 20

True Negatives(TN) = 755

False Positives(FP) = 38

False Negatives(FN) = 1

```
In [157]: sns.heatmap(cm, annot=True)
plt.title('Confusion Matrix')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```



```
In [158]: #SVM with Linear kernel
```

```
linear_svc = SVC(kernel='linear')
```

```
In [159]: linear_svc.fit(X_train_svm, Y_train_svm)
```

```
Out[159]:
```

```
SVC
SVC(kernel='linear')
```

```
In [160]: Y_predLinear_svm = linear_svc.predict(X_test_svm)
```

```
In [161]: acc_svm = accuracy_score(Y_test_svm, Y_predLinear_svm)
print("Accuracy Score with Linear Kernel : {:.2f}%". format(acc_svm*100))
```

Accuracy Score with Linear Kernel : 88.42%

```
In [162]: print(classification_report(Y_test_svm, Y_predLinear_svm))
```

	precision	recall	f1-score	support
Hate Speech Detected	0.40	0.30	0.34	290
NO hate and offensive speech	0.84	0.82	0.83	835
Offensive language detected	0.92	0.94	0.93	3832
accuracy			0.88	4957
macro avg	0.72	0.69	0.70	4957
weighted avg	0.88	0.88	0.88	4957

```
In [163]: cm = confusion_matrix(Y_test_svm, Y_predLinear_svm)
```

```
print('Confusion matrix : ')
print(cm)
print('True Positives(TP) = ', cm[0,0])
print('True Negatives(TN) = ', cm[1,1])
print('False Positives(FP) = ', cm[0,1])
print('False Negatives(FN) = ', cm[1,0])
```

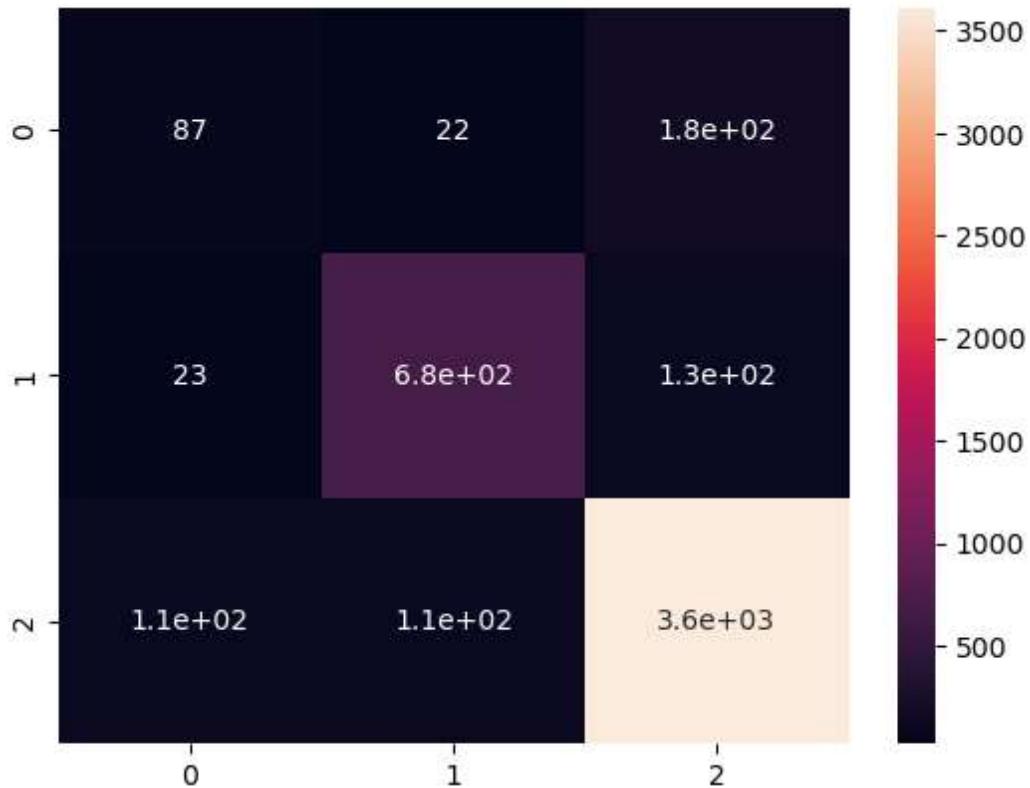
```
Confusion matrix :
```

```
[[ 87  22 181]
 [ 23 684 128]
 [109 111 3612]]
```

```
True Positives(TP) = 87
True Negatives(TN) = 684
False Positives(FP) = 22
False Negatives(FN) = 23
```

In [164]: `sns.heatmap(cm, annot=True)`

Out[164]: <Axes: >



In [165]: `#SVM with RBF Kernel`

```
rbf_svc = SVC(kernel='rbf')
```

In [166]: `rbf_svc.fit(X_train_svm, Y_train_svm)`

Out[166]:

```
▼ SVC
  SVC()
```

In [167]: `Y_predRBF_svm = rbf_svc.predict(X_test_svm)`

In [168]: `acc_rbf_svm = accuracy_score(Y_test_svm, Y_predRBF_svm)
print("Accuracy Score with RBF Kernel : {:.2f}%". format(acc_rbf_svm*100))`

Accuracy Score with RBF Kernel : 89.29%

```
In [169]: print(classification_report(Y_test_svm, Y_predRBF_svm))
```

	precision	recall	f1-score	support
Hate Speech Detected	0.43	0.07	0.12	290
NO hate and offensive speech	0.80	0.90	0.85	835
Offensive language detected	0.92	0.95	0.94	3832
accuracy			0.89	4957
macro avg	0.71	0.64	0.63	4957
weighted avg	0.87	0.89	0.87	4957

```
In [170]: cm = confusion_matrix(Y_test_svm, Y_predRBF_svm)
```

```
print('Confusion matrix : ')
print(cm)
print('True Positives(TP) = ', cm[0,0])
print('True Negatives(TN) = ', cm[1,1])
print('False Positives(FP) = ', cm[0,1])
print('False Negatives(FN) = ', cm[1,0])
```

```
Confusion matrix :
```

```
[[ 20   38  232]
 [  1  755   79]
 [ 26  155 3651]]
```

```
True Positives(TP) = 20
```

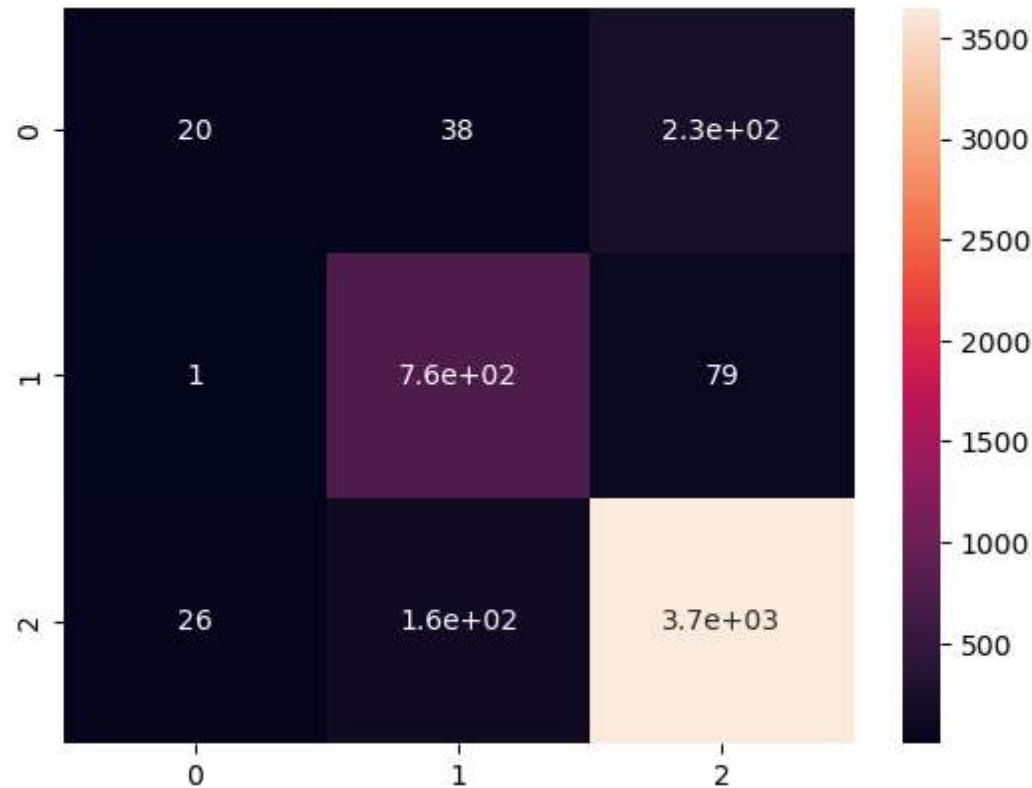
```
True Negatives(TN) = 755
```

```
False Positives(FP) = 38
```

```
False Negatives(FN) = 1
```

```
In [171]: sns.heatmap(cm, annot=True)
```

```
Out[171]: <Axes: >
```



```
In [172]: #SVM with Sigmoid Kernel
```

```
sigmoid_svc = SVC(kernel='sigmoid')
```

```
In [173]: sigmoid_svc.fit(X_train_svm, Y_train_svm)
```

```
Out[173]:
```

```
SVC  
SVC(kernel='sigmoid')
```

```
In [174]: Y_predsigmoid_svm = sigmoid_svc.predict(X_test_svm)
```

```
In [175]: acc_sigmoid_svm = accuracy_score(Y_test_svm, Y_predsigmoid_svm)  
print("Accuracy Score with Sigmoid Kernel : {:.2f}%". format(acc_sigmoid_svm*100))
```

Accuracy Score with Sigmoid Kernel : 88.78%

```
In [176]: print(classification_report(Y_test_svm, Y_predsigmoid_svm))
```

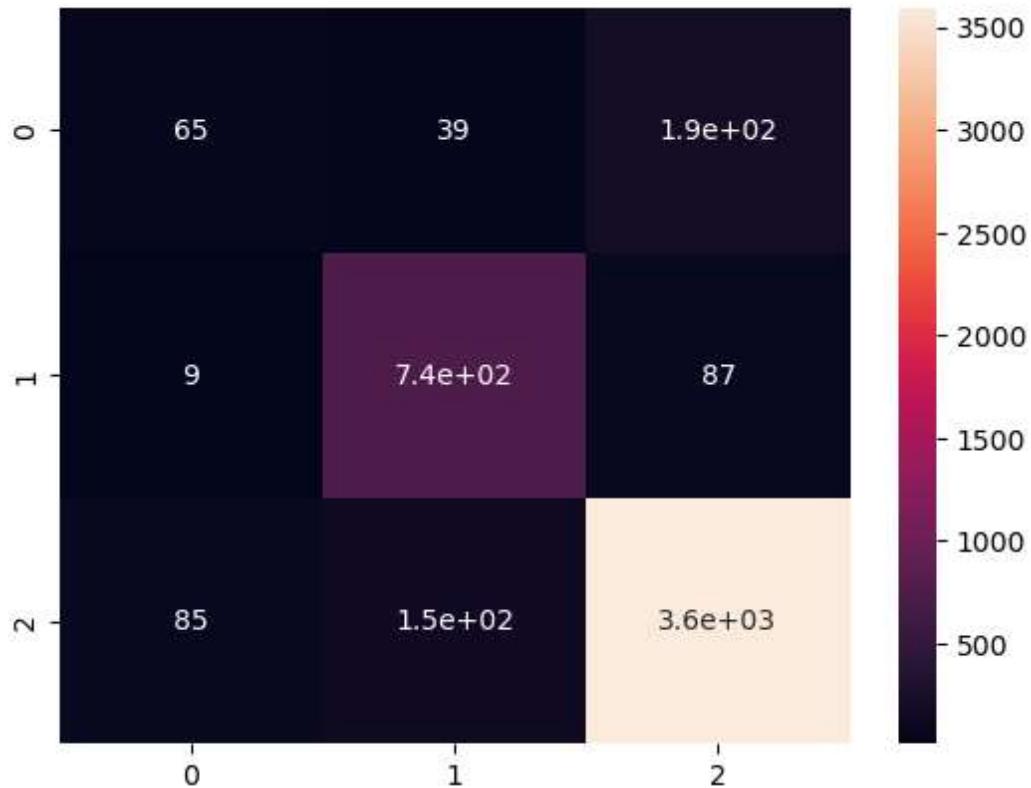
	precision	recall	f1-score	support
Hate Speech Detected	0.41	0.22	0.29	290
NO hate and offensive speech	0.80	0.89	0.84	835
Offensive language detected	0.93	0.94	0.93	3832
accuracy			0.89	4957
macro avg	0.71	0.68	0.69	4957
weighted avg	0.88	0.89	0.88	4957

```
In [177]: cm = confusion_matrix(Y_test_svm, Y_predsigmoid_svm)
print('Confusion matrix : ')
print(cm)
print('True Positives(TP) = ', cm[0,0])
print('True Negatives(TN) = ', cm[1,1])
print('False Positives(FP) = ', cm[0,1])
print('False Negatives(FN) = ', cm[1,0])
```

```
Confusion matrix :
[[ 65   39  186]
 [  9  739   87]
 [ 85  150 3597]]
True Positives(TP) = 65
True Negatives(TN) = 739
False Positives(FP) = 39
False Negatives(FN) = 9
```

```
In [178]: sns.heatmap(cm, annot=True)
```

```
Out[178]: <Axes: >
```



```
In [179]: #SVM with Polynomial Kernel
```

```
poly_svc = SVC(kernel='poly')
```

```
In [180]: poly_svc.fit(X_train_svm, Y_train_svm)
```

```
Out[180]:
```

```
SVC  
SVC(kernel='poly')
```

```
In [181]: Y_predPoly_svm = poly_svc.predict(X_test_svm)
```

```
In [182]: acc_poly_svm = accuracy_score(Y_test_svm, Y_predPoly_svm)  
print("Accuracy Score with Polynomial Kernel : {:.2f}%". format(acc_poly_svm*100))
```

Accuracy Score with Polynomial Kernel : 80.11%

```
In [183]: print(classification_report(Y_test_svm, Y_predPoly_svm))
```

	precision	recall	f1-score	support
Hate Speech Detected	0.31	0.01	0.03	290
NO hate and offensive speech	0.86	0.20	0.32	835
Offensive language detected	0.80	0.99	0.89	3832
accuracy			0.80	4957
macro avg	0.66	0.40	0.41	4957
weighted avg	0.78	0.80	0.74	4957

```
In [184]: cm = confusion_matrix(Y_test_svm, Y_predPoly_svm)
```

```
print('Confusion matrix : ')
print(cm)
print('True Positives(TP) = ', cm[0,0])
print('True Negatives(TN) = ', cm[1,1])
print('False Positives(FP) = ', cm[0,1])
print('False Negatives(FN) = ', cm[1,0])
```

```
Confusion matrix :
```

```
[[ 4   4 282]
 [ 2 164 669]
 [ 7  22 3803]]
```

```
True Positives(TP) = 4
```

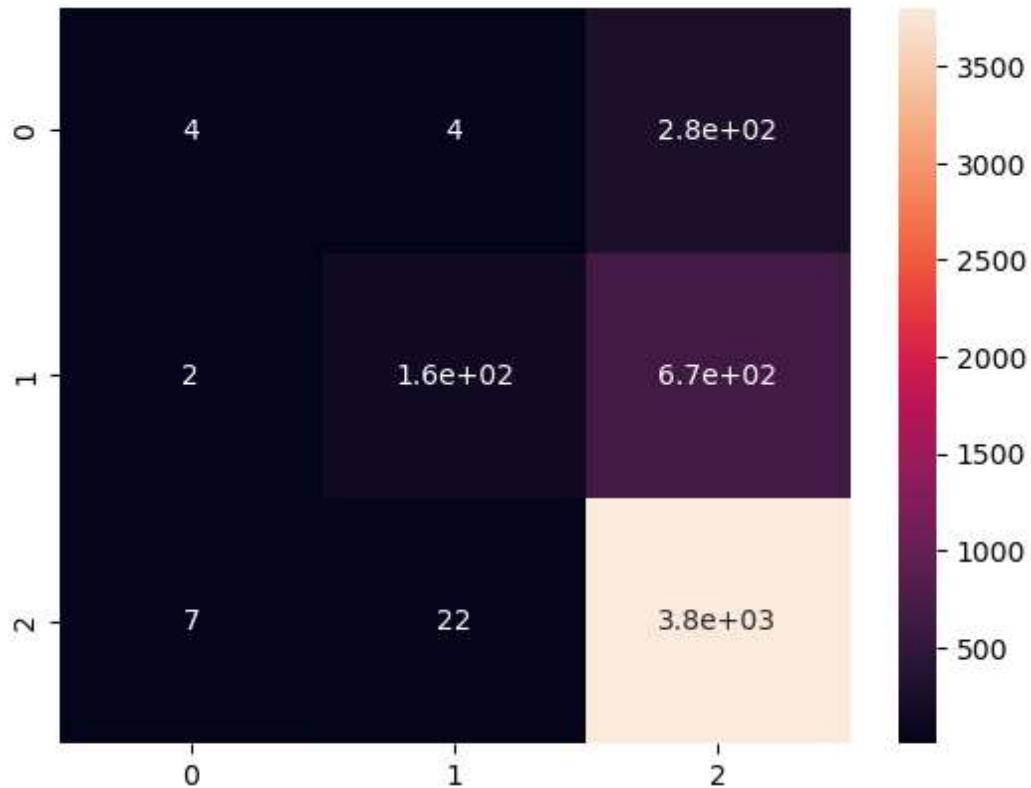
```
True Negatives(TN) = 164
```

```
False Positives(FP) = 4
```

```
False Negatives(FN) = 2
```

In [185]: `sns.heatmap(cm, annot=True)`

Out[185]: <Axes: >



## Tuning SVM

In [186]: `#Grid Search`

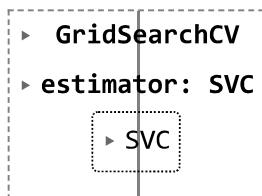
```
param_grid = {
    'C' : [0.1, 1, 10, 100],
    'kernel' : ['linear', 'rbf', 'poly', 'sigmoid']
}
```

In [187]: `svcm = SVC()`

In [188]: `grid_search = GridSearchCV(svcm, param_grid, cv=5)`

In [189]: `grid_search.fit(X_train_svm, Y_train_svm)`

Out[189]:



```
In [190]: best_param_svm = grid_search.best_params_
print("Best hyperparameter : ", best_param_svm)
```

Best hyperparameter : {'C': 0.1, 'kernel': 'linear'}

```
In [191]: best_svm = SVC(C=best_param_svm['C'], kernel=best_param_svm['kernel'])
```

```
In [192]: best_svm.fit(X_train_svm, Y_train_svm)
```

Out[192]:

```
SVC
SVC(C=0.1, kernel='linear')
```

```
In [193]: Y_pred_tun_svm = best_svm.predict(X_test_svm)
```

```
In [194]: acc_tun_svm = accuracy_score(Y_test_svm, Y_pred_tun_svm)
print("Accuracy after Grid Search : {:.2f}%". format(acc_tun_svm * 100))
```

Accuracy after Grid Search : 89.71%

```
In [195]: print(classification_report(Y_test_svm, Y_pred_svm))
```

	precision	recall	f1-score	support
Hate Speech Detected	0.43	0.07	0.12	290
NO hate and offensive speech	0.80	0.90	0.85	835
Offensive language detected	0.92	0.95	0.94	3832
accuracy			0.89	4957
macro avg	0.71	0.64	0.63	4957
weighted avg	0.87	0.89	0.87	4957

```
In [196]: cm = confusion_matrix(Y_test_svm, Y_pred_tun_svm)
print('Confusion matrix : ')
print(cm)
print('True Positives(TP) = ', cm[0,0])
print('True Negatives(TN) = ', cm[1,1])
print('False Positives(FP) = ', cm[0,1])
print('False Negatives(FN) = ', cm[1,0])
```

Confusion matrix :

```
[[ 25  40 225]
 [  6 769  60]
 [ 10 169 3653]]
```

True Positives(TP) = 25

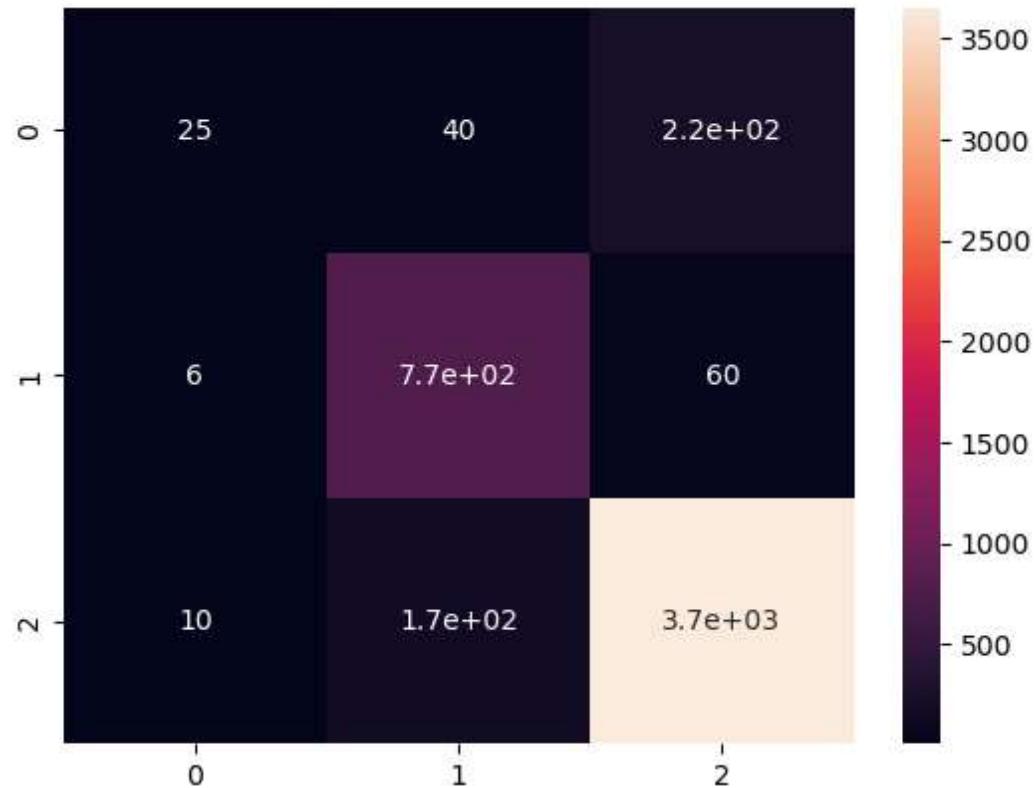
True Negatives(TN) = 769

False Positives(FP) = 40

False Negatives(FN) = 6

```
In [197]: sns.heatmap(cm, annot=True)
```

```
Out[197]: <Axes: >
```



```
In [198]: #Tuning the model with Random Search
```

```
param_grid = {
    'C' : [0.1, 1, 10, 100],
    'kernel' : ['linear', 'rbf', 'poly', 'sigmoid']
}
```

```
In [199]: svcmm = SVC()
```

```
In [200]: from sklearn.model_selection import RandomizedSearchCV
random_search = RandomizedSearchCV(svcmm, param_grid, cv=5)
```

```
In [201]: random_search.fit(X_train_svm, Y_train_svm)
```

```
Out[201]:
```

```
  ▶ RandomizedSearchCV
      ▶ estimator: SVC
          ▶ SVC
```

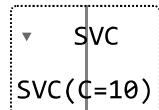
```
In [202]: best_param_rs_svm = random_search.best_params_
print("Best Parameter : ", best_param_rs_svm)
```

Best Parameter : {'kernel': 'rbf', 'C': 10}

```
In [203]: best_svm_rs = SVC(C=best_param_rs_svm['C'], kernel=best_param_rs_svm['kernel'])
```

```
In [204]: best_svm_rs.fit(X_train_svm, Y_train_svm)
```

```
Out[204]:
```



```
In [205]: Y_pred_rs_svm = best_svm_rs.predict(X_test_svm)
```

```
In [206]: acc_rs_svm = accuracy_score(Y_test_svm, Y_pred_rs_svm)
print("Accuracy after Randomized Search: {:.2f}%".format(acc_rs_svm * 100))
```

Accuracy after Randomized Search: 89.51%

```
In [207]: print(classification_report(Y_test_svm, Y_pred_rs_svm))
```

	precision	recall	f1-score	support
Hate Speech Detected	0.48	0.22	0.30	290
NO hate and offensive speech	0.84	0.86	0.85	835
Offensive language detected	0.92	0.95	0.94	3832
accuracy			0.90	4957
macro avg	0.75	0.68	0.69	4957
weighted avg	0.88	0.90	0.89	4957

```
In [208]: cm = confusion_matrix(Y_test_svm, Y_pred_rs_svm)
print('Confusion matrix : ')
print(cm)
print('True Positives(TP) = ', cm[0,0])
print('True Negatives(TN) = ', cm[1,1])
print('False Positives(FP) = ', cm[0,1])
print('False Negatives(FN) = ', cm[1,0])
```

Confusion matrix :

```
[[ 63  27 200]
 [  8 715 112]
 [ 61 112 3659]]
```

True Positives(TP) = 63

True Negatives(TN) = 715

False Positives(FP) = 27

False Negatives(FN) = 8

```
In [209]: sns.heatmap(cm, annot=True)
```

```
Out[209]: <Axes: >
```



## Logistic Regression

```
In [322]: from sklearn.linear_model import LogisticRegression
```

```
In [323]: x_lor=np.array(df["tweet"])
y_lor=np.array(df["labels"])
```

```
In [324]: cv_lor=CountVectorizer()
x_lor=cv_lor.fit_transform(x_lor)
```

```
In [325]: X_train_lor, X_test_lor, Y_train_lor, Y_test_lor = train_test_split(x_lor, y_lor,
```

```
In [326]: lor = LogisticRegression()

lor.fit(X_train_lor,Y_train_lor)
y_predicted_lor = lor.predict(X_test_lor)
```

```
In [327]: lor.predict_proba(X_test_lor)
```

```
Out[327]: array([[4.67877529e-03, 2.55407708e-04, 9.95065817e-01],
 [1.52501681e-03, 8.00522536e-06, 9.98466978e-01],
 [2.22506802e-03, 1.20128779e-01, 8.77646153e-01],
 ...,
 [9.90904192e-03, 2.74413957e-04, 9.89816544e-01],
 [1.68439435e-01, 1.05232980e-02, 8.21037268e-01],
 [1.12164226e-02, 1.47285981e-03, 9.87310718e-01]])
```

```
In [328]: lor.score(X_test_lor, Y_test_lor)
```

```
Out[328]: 0.8926770223925762
```

```
In [329]: lor.coef_
```

```
Out[329]: array([[ 0.28616955, 0.          , -0.08501712, ... , -0.01417363,
 -0.00780366, 0.42725769],
 [-0.06017044, 0.          , -0.00061859, ... , -0.00194629,
 0.13422793, -0.02672619],
 [-0.22599911, 0.          , 0.08563571, ... , 0.01611992,
 -0.12642428, -0.4005315 ]])
```

```
In [330]: lor.intercept_ #y_intercept
```

```
Out[330]: array([-1.75531072, 0.78914823, 0.96616249])
```

```
In [331]: import math
```

```
def sigmoid(x):
    return 1/(1 + math.exp(-x))
```

```
In [332]: def prediction_function(age):
    z= 0.14915498 * age - 5.80219736 # y = B0 + B1 * x
    y = sigmoid(z)
    return y
```

```
In [333]: #Calculating Test Score
print('Test Score : ', lor.score(X_test_lor, Y_test_lor))
```

```
Test Score : 0.8926770223925762
```

```
In [334]: #Predicting the train value
pred_train = lor.predict(X_train_lor)
```

```
In [335]: #Predicting the test value
pred_test = lor.predict(X_test_lor)
```

```
In [336]: #Classification report for the Training set
print(classification_report(Y_train_lor, pred_train))
```

	precision	recall	f1-score	support
Hate Speech Detected	0.97	0.66	0.79	1140
NO hate and offensive speech	0.97	0.98	0.98	3328
Offensive language detected	0.97	1.00	0.98	15358
accuracy			0.97	19826
macro avg	0.97	0.88	0.92	19826
weighted avg	0.97	0.97	0.97	19826

## Tuning - LOR

```
In [225]: param_grid = {
    'penalty' : ['l1', 'l2'],
    'C' : [0.1, 0.5, 1, 5, 10]
}
```

#l1 Lasso L2 Ridge

```
In [226]: grid = GridSearchCV(estimator=lor, param_grid=param_grid, cv=5)
```

```
In [227]: grid.fit(X_train_lor,Y_train_lor)
```

Out[227]:

```
▶      GridSearchCV
  ▶ estimator: LogisticRegression
    ▶ LogisticRegression
```

```
In [228]: #Getting the best hyperparameters
best_params_lor = grid.best_params_

# Getting the best model
best_lor = grid.best_estimator_
```

```
In [229]: # Making predictions on the test set
Y_pred_tun_lor = best_lor.predict(X_test_lor)
```

```
In [230]: # Evaluating the best model
acc_tun_lor = accuracy_score(Y_test_lor, Y_pred_tun_lor)
print("Accuracy after Grid Search : {:.2f}%". format(acc_tun_lor * 100))
```

Accuracy after Grid Search : 89.57%

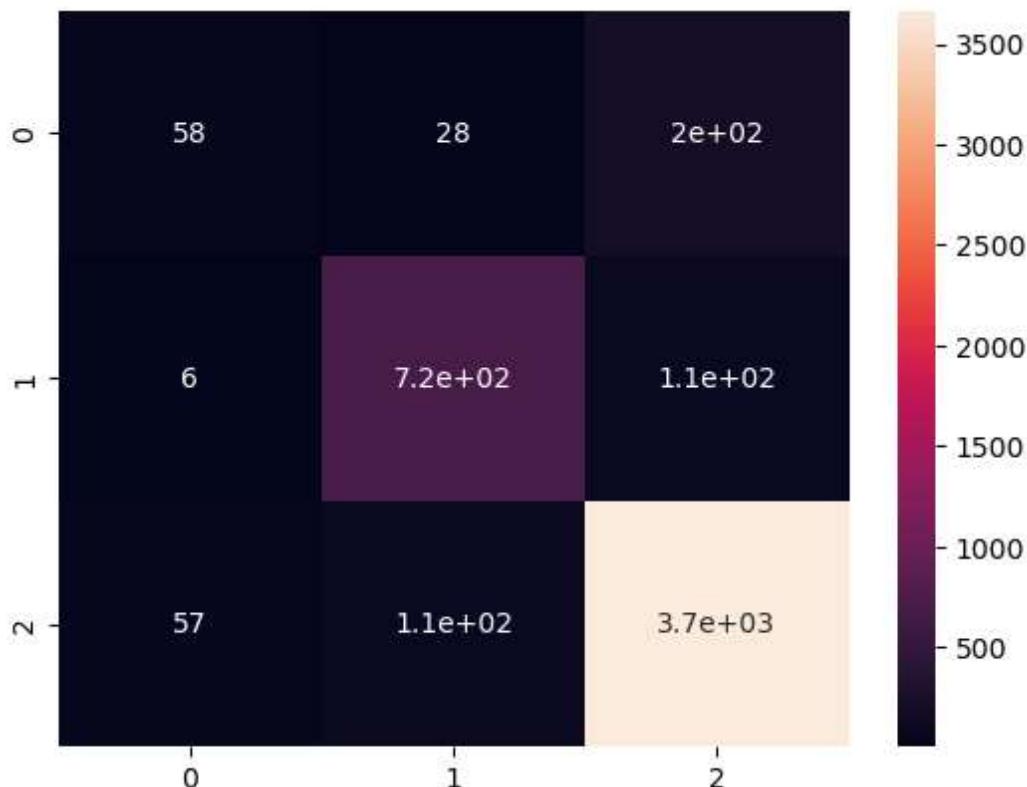
```
In [231]: #Confusion Matrix after Grid Search
confusion = confusion_matrix(Y_test_lor, Y_pred_tun_lor)
print('Confusion Matrix:')
print(confusion)
```

Confusion Matrix:

```
[[ 58   28  204]
 [  6  715  114]
 [ 57  108 3667]]
```

```
In [232]: sns.heatmap(confusion, annot=True)
```

Out[232]: <Axes: >



```
In [233]: #Classification report after Grid Search
classification_rep = classification_report(Y_test_lor, Y_pred_tun_lor)
print('Classification Report:')
print(classification_rep)
```

Classification Report:

	precision	recall	f1-score	support
Hate Speech Detected	0.48	0.20	0.28	290
NO hate and offensive speech	0.84	0.86	0.85	835
Offensive language detected	0.92	0.96	0.94	3832
accuracy			0.90	4957
macro avg	0.75	0.67	0.69	4957
weighted avg	0.88	0.90	0.88	4957

```
In [234]: from scipy.stats import uniform

# Defining hyperparameters and their possible values for tuning
param_dist = {
    'penalty': ['l1', 'l2'],
    'C': uniform(loc=0, scale=4) # Random values between 0 and 4 for regularization
}
```

```
In [235]: # Creating a RandomizedSearchCV object
random_search = RandomizedSearchCV(lor, param_distributions=param_dist, n_iter=10)
```

```
In [236]: # Fitting the RandomizedSearchCV object to the training data
random_search.fit(X_train_lor, Y_train_lor)
```

Out[236]:

```
RandomizedSearchCV
  estimator: LogisticRegression
    LogisticRegression
```

```
In [237]: # Getting the best hyperparameters
best_params_rs_tun_lor = random_search.best_params_
```

```
In [238]: # Getting the best model
best_lor_rs = random_search.best_estimator_
```

```
In [239]: # Making predictions on the test set
Y_pred_rs_tun_lor = best_lor_rs.predict(X_test_lor)
```

```
In [240]: # Evaluating the best model
acc_rs_tun_lor = accuracy_score(Y_test_lor, Y_pred_rs_tun_lor)
acc_rs_tun_lor
```

Out[240]: 0.8961065160379261

```
In [241]: #Classification report after Randomized Search
classification_rep = classification_report(Y_test_lor, Y_pred_rs_tun_lor)
print('Classification Report:')
print(classification_rep)
```

Classification Report:

	precision	recall	f1-score	support
Hate Speech Detected	0.49	0.20	0.28	290
NO hate and offensive speech	0.84	0.86	0.85	835
Offensive language detected	0.92	0.96	0.94	3832
accuracy			0.90	4957
macro avg	0.75	0.67	0.69	4957
weighted avg	0.88	0.90	0.88	4957

```
In [242]: #Confusion Matrix
confusion = confusion_matrix(Y_test_lor, Y_pred_rs_tun_lor)
print('Confusion Matrix:')
print(confusion)
```

Confusion Matrix:

```
[[ 58   27 205]
 [  5 715 115]
 [ 55 108 3669]]
```

```
In [243]: sns.heatmap(confusion, annot=True)
```

Out[243]: <Axes: >



# Random Forest

```
In [244]: from sklearn.ensemble import RandomForestClassifier
```

```
In [245]: x_rf=np.array(df["tweet"])
y_rf=np.array(df["labels"])
```

```
In [246]: cv_rf=CountVectorizer()
x_rf= cv_rf.fit_transform(x_rf)
```

```
In [247]: X_train_rf, X_test_rf, Y_train_rf, Y_test_rf = train_test_split(x_rf, y_rf, test_
```

```
In [248]: # Create a Random Forest classifier with 100 trees
rf = RandomForestClassifier(n_estimators=100, random_state=42)
```

```
In [249]: # Fit the model on the training data
rf.fit(X_train_rf, Y_train_rf)
```

```
Out[249]:
```

```
RandomForestClassifier
RandomForestClassifier(random_state=42)
```

```
In [250]: # Predict on the test data
y_pred_rf = rf.predict(X_test_rf)
```

```
In [251]: # Calculate accuracy
acc_rf = rf.score(X_test_rf, Y_test_rf)
print("Accuracy:", acc_rf)
```

Accuracy: 0.8916683477910026

```
In [252]: #Classification report
classification_rep_rf = classification_report(Y_test_rf, y_pred_rf)
print('Classification Report:')
print(classification_rep_rf)
```

Classification Report:

	precision	recall	f1-score	support
Hate Speech Detected	0.54	0.27	0.36	290
NO hate and offensive speech	0.82	0.82	0.82	835
Offensive language detected	0.92	0.95	0.94	3832
accuracy			0.89	4957
macro avg	0.76	0.68	0.71	4957
weighted avg	0.88	0.89	0.88	4957

## Tuning - RF

```
In [253]: param_grid = {
    'n_estimators': [100, 200, 300, 400, 500],
    'max_depth': [5, 10, 15, 20, 25],
    'max_features': ['auto', 'sqrt', 'log2'],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
}
```

```
In [254]: rf = RandomForestClassifier(random_state=42)
```

```
In [255]: # Create a RandomizedSearchCV object
random_search = RandomizedSearchCV(rf, param_grid, n_iter=20, cv=5, random_state=42)
```

```
In [256]: # Fit the model to the training data
random_search.fit(X_train_rf, Y_train_rf)
```

```
Out[256]:
```

- ▶ RandomizedSearchCV
- ▶ estimator: RandomForestClassifier
  - ▶ RandomForestClassifier

```
In [257]: # Print the best hyperparameters found
print("Best hyperparameters:", random_search.best_params_)
```

```
Best hyperparameters: {'n_estimators': 200, 'min_samples_split': 2, 'min_samples_leaf': 1, 'max_features': 'auto', 'max_depth': 25}
```

```
In [258]: # Get the best model
best_model = random_search.best_estimator_
```

```
In [259]: # Make predictions on the test set
y_pred_tun_rf = best_model.predict(X_test_rf)
```

```
In [260]: # Evaluate performance
acc_tun_rf = best_model.score(X_test_rf, Y_test_rf)
print("Accuracy:", acc_tun_rf)
```

```
Accuracy: 0.7740568892475288
```

```
In [261]: #Classification report after Randomized Search
classification_rep_tun_rf = classification_report(Y_test_rf, y_pred_tun_rf)
print('Classification Report:')
print(classification_rep_tun_rf)
```

Classification Report:

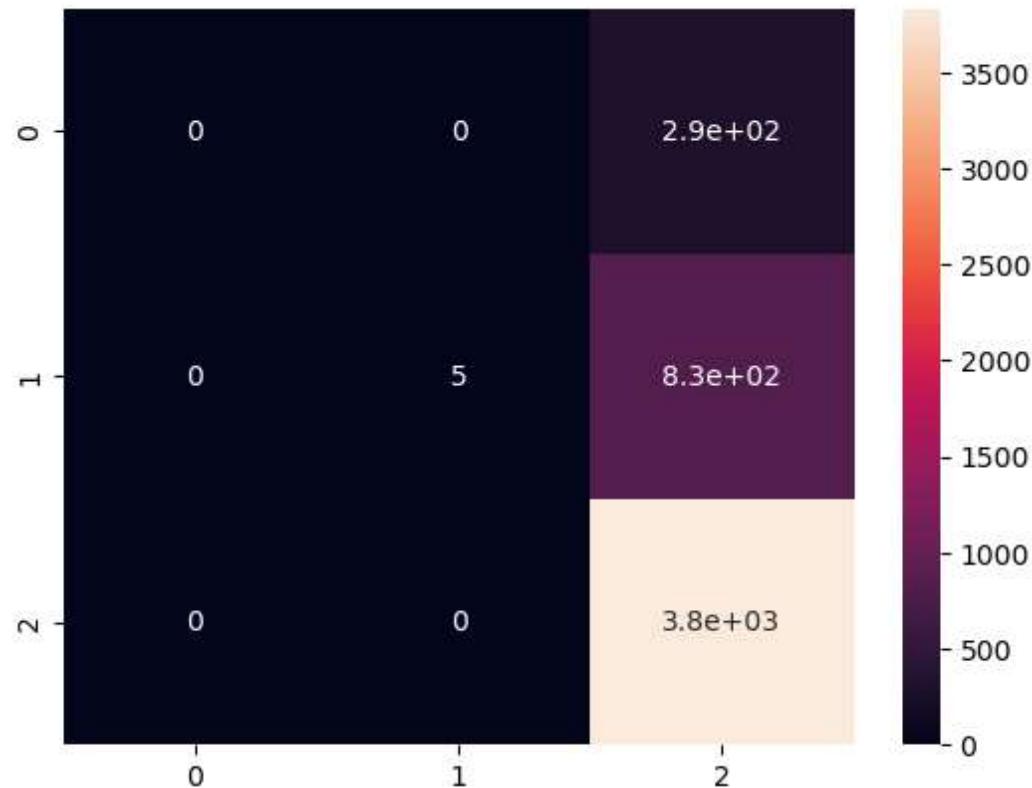
	precision	recall	f1-score	support
Hate Speech Detected	0.00	0.00	0.00	290
NO hate and offensive speech	1.00	0.01	0.01	835
Offensive language detected	0.77	1.00	0.87	3832
accuracy			0.77	4957
macro avg	0.59	0.34	0.29	4957
weighted avg	0.77	0.77	0.68	4957

```
In [262]: confusion = confusion_matrix(Y_test_rf, y_pred_tun_rf)
print('Confusion Matrix:')
print(confusion)
```

Confusion Matrix:  
[[ 0 0 290]  
 [ 0 5 830]  
 [ 0 0 3832]]

```
In [263]: sns.heatmap(confusion, annot=True)
```

Out[263]: <Axes: >



# Naive Bayes

```
In [275]: from sklearn import model_selection, naive_bayes, svm, metrics, feature_extraction
```

```
In [265]: x_nb=np.array(df["tweet"])
y_nb=np.array(df["labels"])
```

```
In [266]: cv_nb=CountVectorizer()
x_nb= cv_nb.fit_transform(x_nb)
```

```
In [267]: X_train_nb, X_test_nb, Y_train_nb, Y_test_nb = train_test_split(x_nb, y_nb, test_
```

```
In [280]: from sklearn.preprocessing import MaxAbsScaler

scaler = MaxAbsScaler()
scaled_data = scaler.fit_transform(X_train_nb)
X_train_nb = scaler.fit_transform(X_train_nb)
X_test_nb = scaler.transform(X_test_nb)
```

```
In [282]: #choosing the naive bayes
bayes = naive_bayes.BernoulliNB()
bayes.fit(X_train_nb,Y_train_nb)
```

```
Out[282]:
```

▼ BernoulliNB  
BernoulliNB()

```
In [283]: # predicting the y value based on bayes model
y_pred_nb=bayes.predict(X_test_nb)
```

```
In [310]: #calculating the accuracy of the model on the testing data
acc_nb=metrics.accuracy_score(Y_test_nb,y_pred_nb)
acc_nb
```

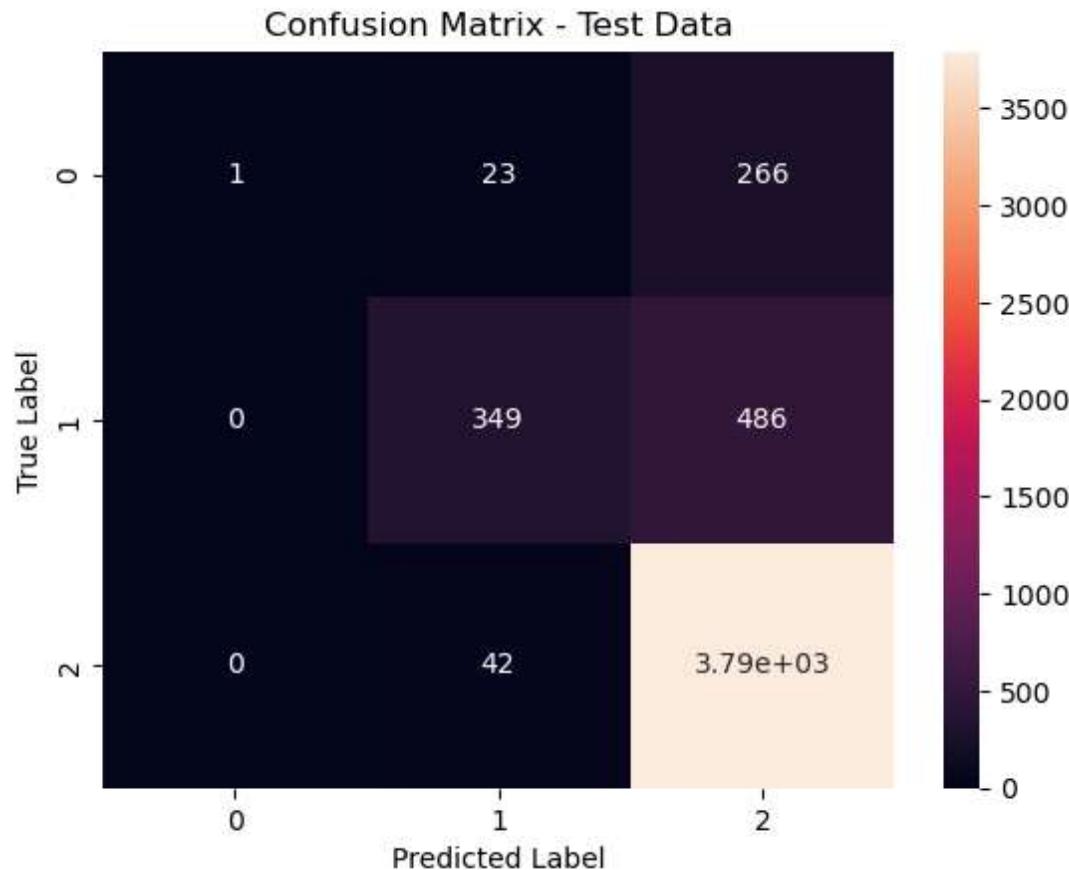
```
Out[310]: 0.8351825701028848
```

```
In [285]: # generating the classification report for the performance measures.
print(metrics.classification_report(Y_test_nb, y_pred_nb))
```

	precision	recall	f1-score	support
Hate Speech Detected	1.00	0.00	0.01	290
NO hate and offensive speech	0.84	0.42	0.56	835
Offensive language detected	0.83	0.99	0.91	3832
accuracy			0.84	4957
macro avg	0.89	0.47	0.49	4957
weighted avg	0.85	0.84	0.79	4957

```
In [286]: cm=confusion_matrix(Y_test_nb,y_pred_nb)
print(cm)
sns.heatmap(cm, annot=True, fmt='.3g')
plt.title('Confusion Matrix - Test Data')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```

```
[[ 1  23 266]
 [ 0 349 486]
 [ 0  42 3790]]
```



## Tuning Naive bayes

```
In [293]: param_grid = {
    'alpha': [0.1, 1, 10, 100],
    'fit_prior': [True, False]
}
```

```
In [294]: bayes = naive_bayes.BernoulliNB()
grid_search = GridSearchCV(bayes, param_grid, cv=5)
grid_search.fit(X_train_nb, Y_train_nb)
```

Out[294]:

```
GridSearchCV
  estimator: BernoulliNB
    BernoulliNB
```

```
In [295]: best_param = grid_search.best_params_
best_nb = naive_bayes.BernoulliNB(alpha = best_param['alpha'], fit_prior = best_p
best_nb.fit(X_train_nb, Y_train_nb)
y_pred_1 = best_nb.predict(X_test_nb)
```

```
In [296]: print("Best Hyperparameter : ", best_param)
```

Best Hyperparameter : {'alpha': 0.1, 'fit\_prior': True}

```
In [308]: acc_nb_tun = accuracy_score(Y_test_nb, y_pred_1)
acc_nb_tun
```

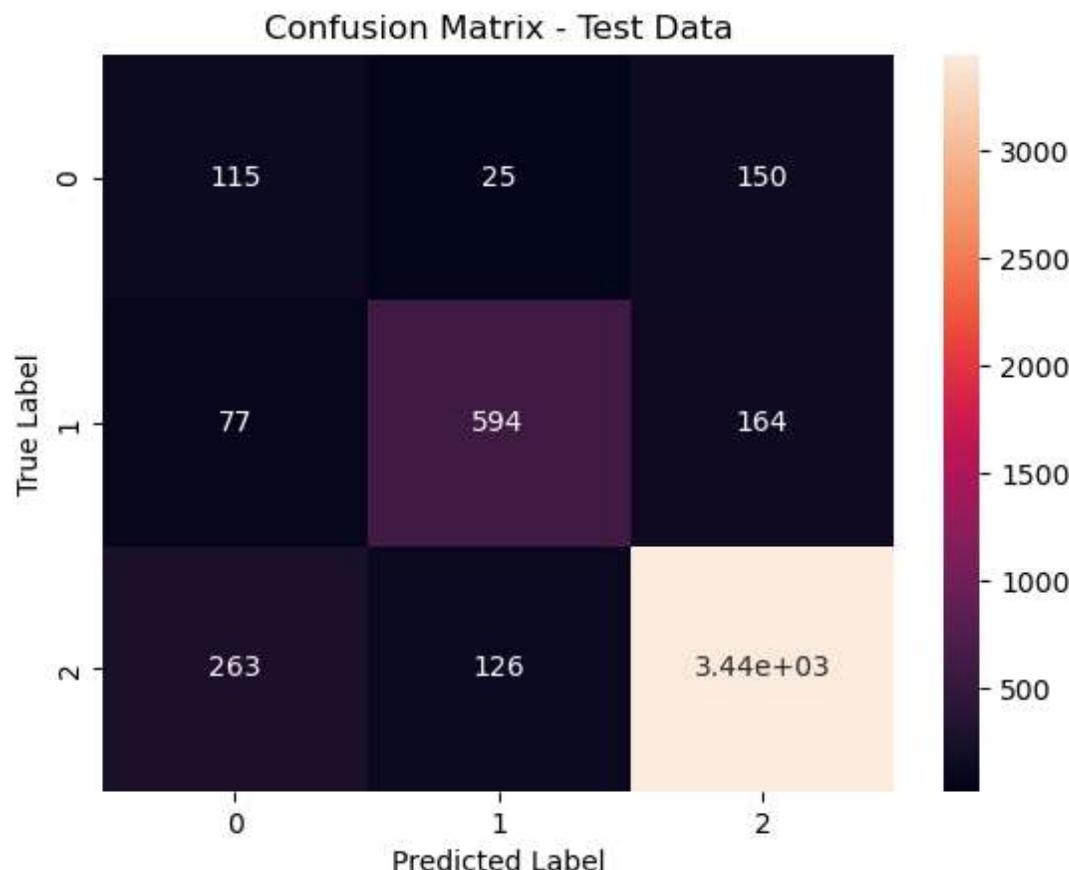
Out[308]: 0.8376033891466613

```
In [309]: print (classification_report(Y_test_nb,y_pred_1))
```

	precision	recall	f1-score	support
Hate Speech Detected	0.25	0.40	0.31	290
NO hate and offensive speech	0.80	0.71	0.75	835
Offensive language detected	0.92	0.90	0.91	3832
accuracy			0.84	4957
macro avg	0.66	0.67	0.66	4957
weighted avg	0.86	0.84	0.85	4957

```
In [299]: cm=confusion_matrix(Y_test_nb,y_pred_1)
print(cm)
sns.heatmap(cm, annot=True, fmt='.3g')
plt.title('Confusion Matrix - Test Data')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```

```
[[ 115   25  150]
 [  77  594  164]
 [ 263  126 3443]]
```



```
In [300]: #Randomized Search
param_dist = {
    'alpha': uniform(0.1, 2.0), # Example: Uniform distribution for alpha
    'fit_prior':[True,False]
}
```

```
In [304]: bayes = naive_bayes.BernoulliNB()
x_new_nb=scaler.fit_transform(x_nb)
```

```
In [305]: randomized_search = RandomizedSearchCV(bayes, param_distributions=param_dist, n_iter=10)
randomized_search.fit(x_new_nb, y_nb)
```

Out[305]:

```
▶ RandomizedSearchCV
  ▶ estimator: BernoulliNB
    ▶ BernoulliNB
```

```
In [306]: best_param = randomized_search.best_params_
print("Best Hyperparameter : ", best_param)
```

Best Hyperparameter : {'alpha': 0.5152534796161073, 'fit\_prior': False}

```
In [307]: best_nb = naive_bayes.BernoulliNB(alpha = best_param['alpha'], fit_prior = best_param['fit_prior'])
best_nb.fit(X_train_nb, Y_train_nb)
y_pred_2 = best_nb.predict(X_test_nb)
```

```
In [311]: acc_nb_tun_rs = accuracy_score(Y_test_nb, y_pred_2)
acc_nb_tun_rs
```

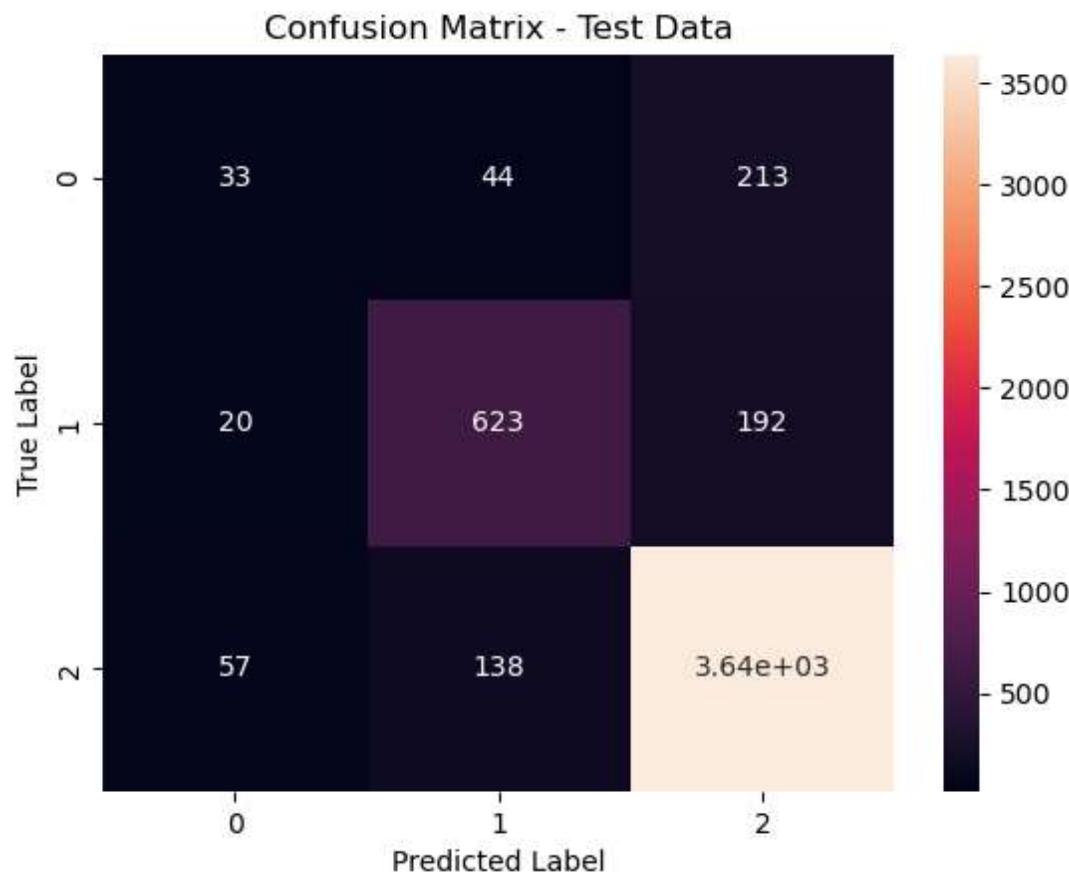
Out[311]: 0.866048012911035

```
In [312]: print (classification_report(Y_test_nb,y_pred_2))
```

	precision	recall	f1-score	support
Hate Speech Detected	0.30	0.11	0.17	290
NO hate and offensive speech	0.77	0.75	0.76	835
Offensive language detected	0.90	0.95	0.92	3832
accuracy			0.87	4957
macro avg	0.66	0.60	0.62	4957
weighted avg	0.84	0.87	0.85	4957

```
In [313]: cm=confusion_matrix(Y_test_nb,y_pred_2)
print(cm)
sns.heatmap(cm, annot=True, fmt='.3g')
plt.title('Confusion Matrix - Test Data')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```

```
[[ 33   44  213]
 [ 20  623  192]
 [ 57  138 3637]]
```



## Pickle

```
In [316]: import pickle as pkl
```

```
In [317]: pkl.dump(svc,open('model.pkl','wb'))
```

```
In [ ]:
```

